



Every industry's best agents will be trained in our worlds

An AI-native neolab building environments, verifiers
and data for self-improving AI

Founded

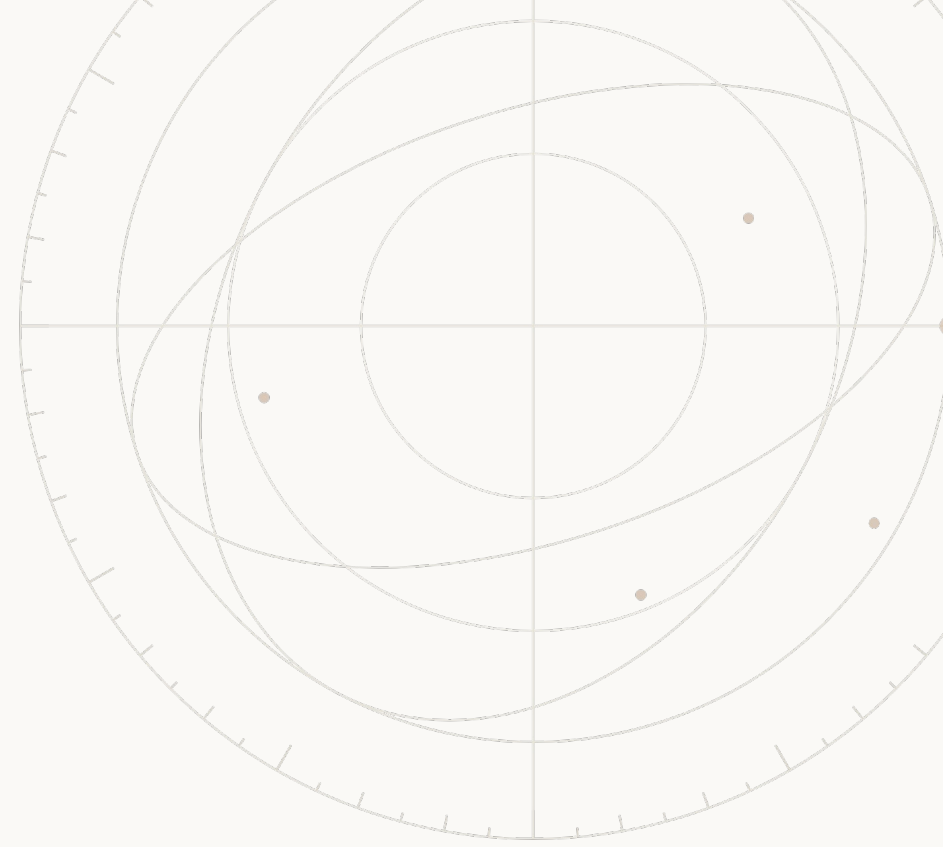
September 10, 2026

Raising

\$6M seed round

Prepared for

Optiver PSI



Overview

An AI-native neolab building the environments, verifiers and data that let agents improve themselves

Revenue

\$3M ARR

Reached 24 days after founding
No outside capital

Flagship

Trading environment

Market making, options, baskets;
risk-adjusted P&L on a held-out day

Early result

+12%

On held-out trading days, after training
in our trading environment (best of several runs)

Team

Three founders, all 21

Cambridge, LSE and Duke mathematics
Jane Street, Citadel, Optiver, Millennium

Open source

600+ GitHub stars

Across our open benchmarks;
7 environments built in 14 days

Market

\$8.5B a year

Spent today on training data
and RL environments

Team

Three 21-year-old founders. Together they turned down \$1M+ in first-year pay to build SimReal



Charles CEO · 21

Employee #1 at United Stables: \$U grew from zero to \$1.5B+ and listed on Binance within a month

Led institutional partnerships with YZi Labs, SIG, DRW and Wintermute

Only intern on HSBC's HKD stablecoin project, working on HKMA compliance

Hedge fund intern, Citadel

LSE, Mathematics
USAMO qualifier

Henry CTO · 21

Cambridge Mathematics, St John's College, First Class with scholarship

Youngest researcher at the Cambridge AI research center, under Prof. Po-Ling Loh (IMS Fellow)

Quant experience at Jane Street, Citadel and Optiver

Top 30 worldwide, Cambridge STEP maths exam
British Physics Olympiad Top Gold

Amaris COO · 21

Data scientist, Millennium Hong Kong

First graduate hire on Millennium's alternative-data team

Co-organized Plug and Play's first Hong Kong event, with HKSTP (200+ attendees)

Duke, Mathematics and Statistics

Problem

Agents can reason. **They still fail at real work**

	Looks right		What actually happens
Trading	Profitable in backtest	→	Loses money live
Software	All tests pass	→	Breaks in production
Accounting	Books look closed	→	Month-end doesn't reconcile
Forecasting	Analysis sounds right	→	Outcome proves it wrong

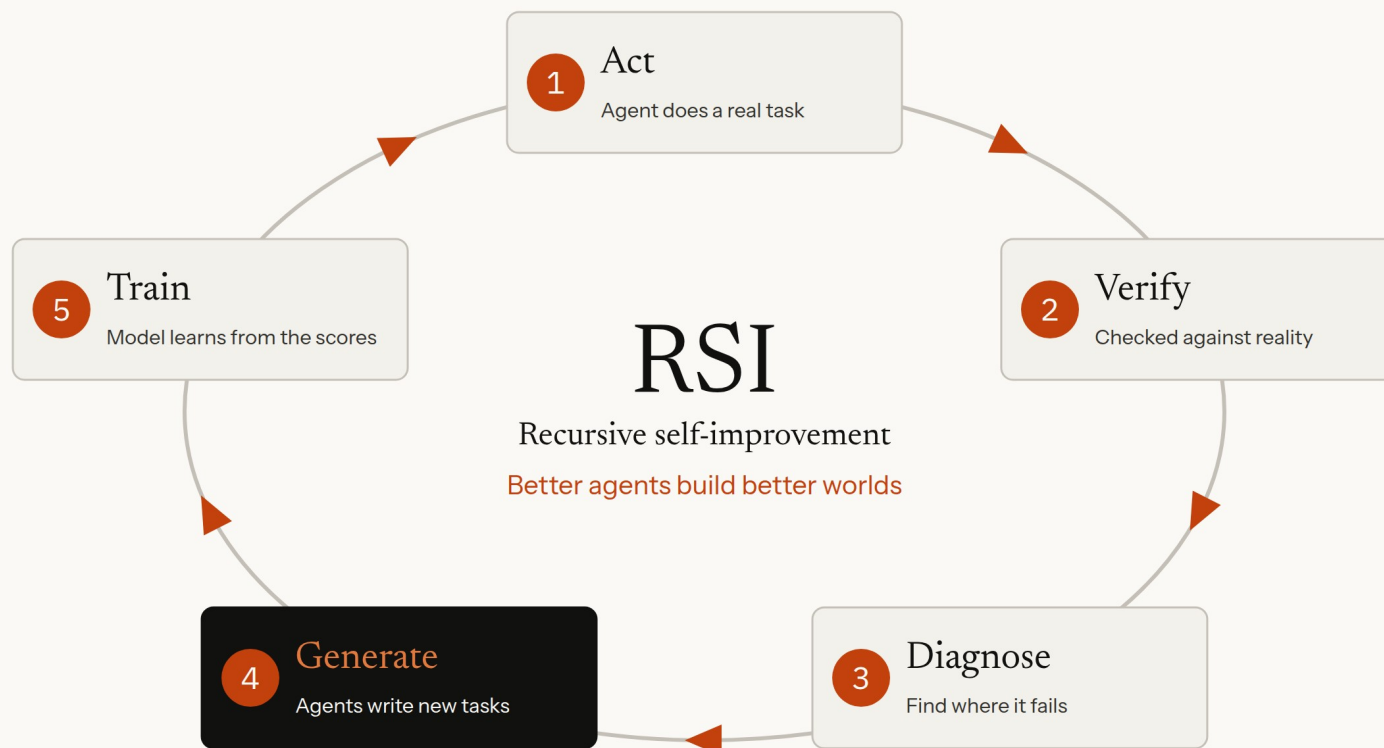
6/32

In a live trading contest (Alpha Arena), frontier models made money in only 6 of 32 runs

The gap: **agents have nowhere realistic to practice, and no source of truth to learn from**

Solution

A loop that turns real outcomes into better agents, **with AI inside it from day one**



What's different

Today

Tasks, labels and rubrics are mostly written by hand

SimReal

Agents help write new tasks and data, find failure modes and improve the environments they train in. Next, they will write the verifiers too. AI-native from the start.

Early result

+12%

on trading days it never saw,
after training in our loop

+12% is the best of several independent replications, measured on trading days held out from training.

Product

Trading is the flagship environment. **The same engine runs ML research, forecasting and more**

FLAGSHIP

•

TRADING

Xitadel

open source

An agent researches a market, codes a strategy and is scored on a held-out trading day it never saw.

Seven published tasks, five types

Single-product market making

Options vs. underlying

Baskets vs. components

Conversion with tariffs and storage

Multi-asset, disclosed counterparties

+12%

on held-out trading days after RL training (best of several runs)

Scored on risk-adjusted P&L; best human strategy = 80

Supporting environments

ML research

SimReal-MLBench

60 past Kaggle competitions, ranked on the final human leaderboard

Forecasting

FuturePredict

Questions with no answer yet, scored on the resolved outcome against a sealed baseline

Also built

Math proofs (MathmoBench) · Reasoning (Puzzle Benchmark)
Software (SWE-Forward) · Accounting (Month-End Close)

600+ GitHub stars across our open benchmarks

SWE-Forward and Month-End Close are private. Best human: the best real strategy on each task from IMC Prosperity, a trading competition.

Inside the trading environment

Research, code, then a private held-out day. **Scored on risk-adjusted P&L against the best human**

01

Research

Historical order books and trades for 2-8 training days, with position limits and product rules



02

Code

The agent backtests locally and submits a Python strategy: `Trader.run(state) → orders`



03

Held-out replay

Run step by step on a private test day. At each step the strategy sees only the order books, reported trades and its own position

Published results · score out of 100, best human = 80

Best human	80
GPT 6	77.3
GLM 5.3	30.1
Kimi K3	27.7
DeepSeek V4 Pro	24.6

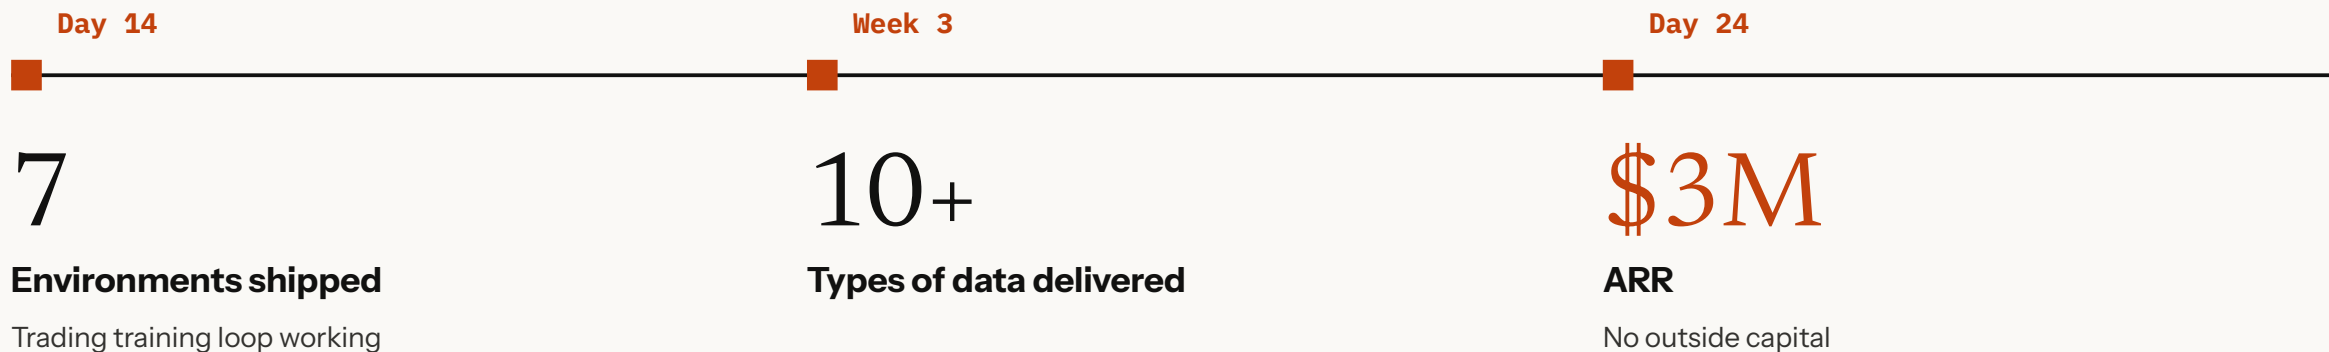
Best human: the best real strategy on each task from IMC Prosperity, a trading competition. Seven published tasks across five task types.

Score and controls

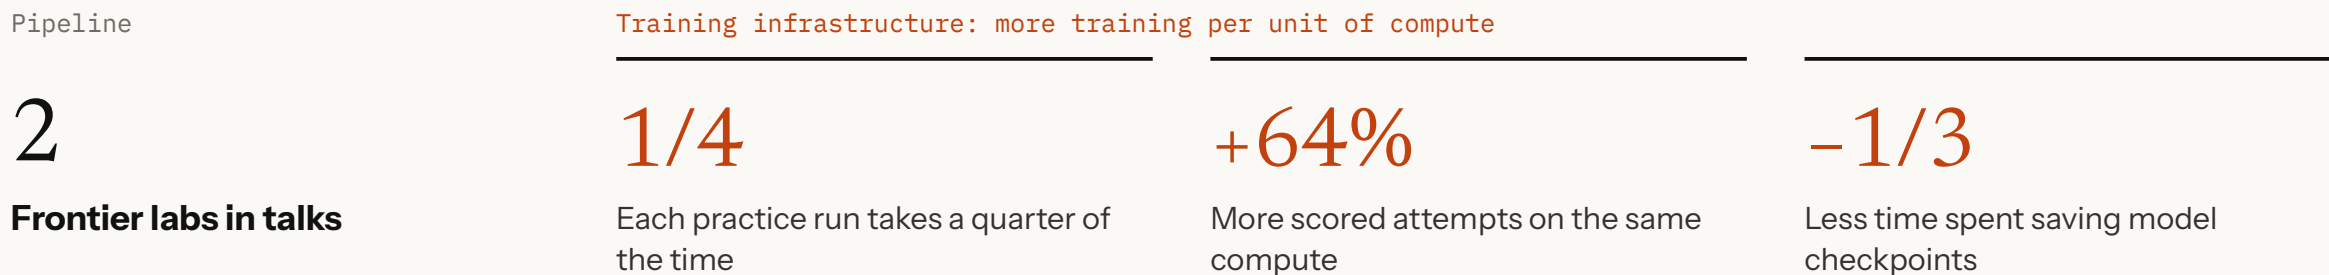
Score	(P&L - 0.1 × max drawdown) × Sharpe factor (0.8-1.0) Rescaled so the best human = 80 and twice it = 100
Selection	Tasks fixed before any model was scored; reserve tasks replace test days once they are no longer hidden
Audit	Full results hash-committed (SHA-256); the training reward is kept separate from this rubric

Traction

Founded September 10, 2026. **\$3M ARR in 24 days, with no outside capital**



Days counted from founding (September 10, 2026)



Infrastructure figures come from separate internal tests and should not be added together.

Why now

Agents are starting to improve themselves. **Environments are now the bottleneck**

What changed in 2026		What we have built
AI is doing AI research OpenAI: agents do 3.1 days of work for every researcher-day OpenAI, Sep 2026	→	SimReal-MLBench: an exam for AI research 60 past Kaggle competitions, scored against humans
AI is improving itself An agent rewrote its own code and improved 7 times in 8 days Weco AIDE ² , Sep 2026	→	Self-improvement working in trading +12% on held-out trading days (best of several runs)
Environments are the bottleneck When the environment gets noisy, top agents drop from 83.9% to 57.6% Breaking the Environment Wall, Sep 2026	→	Seven environments in 14 days 200K+ experts to build more

Our position

Environments gate self-improvement; they are our core research. **Most neolabs have no revenue; we hit \$3M ARR in 24 days**

Sources: OpenAI (Sep 6, 2026); Weco AI, arXiv 2609.26457; Breaking the Environment Wall, arXiv 2609.29773; Deedy Das neolab list (May 2026). +12% is the best of several replications.

Market

\$8.5B a year today. **One vendor, Mercor, grew revenue 27x in 16 months**

Today: training data and RL environments

\$8.5B / year

Combined revenue of 50+ vendors

2030 scenario: our estimate

→ **\$700B / year**

If AI becomes a \$7T economy and 10% goes to training

27x

Mercor run-rate revenue:
\$75M to \$2B in 16 months

→ **Demand for expert data is growing fast**

10x

Mercor valuation:
\$2B to \$20B in 17 months (in talks)

→ **Investors are paying up**

18x

Snorkel AI's new data business:
growth in one year

→ **Demand is shifting to experts and environments**

Sources: Deedy Das market map (Jul 2026); Mercor (TechCrunch, Dealroom); Snorkel AI (Sep 2026). The 2030 figure is our estimate.

Business model

We sell environments, data and evals. **Frontier labs pay today; enterprise and personal agents next**

Now

Frontier AI labs

They buy

Environments, verified tasks and evals to post-train their models

How we charge

Per environment and per task, on quarterly contracts

Estimated price

\$300K–\$1M+

Per lab, per quarter

Market rates: \$200–\$2,000 a task; about \$300K for a complex environment

Next

Enterprise agents

They buy

Custom environments and tests that prove an agent can do the job

How we charge

A paid pilot, then an annual license

Estimated price

\$200K–\$800K

Per customer, per year

Pilot: \$50K–\$150K over 8–12 weeks

Later

Personal agents

They buy

Fit an agent to one person's habits and strategies before it acts for them

How we charge

A monthly subscription, direct or through agent apps

Estimated price

\$5–\$10

Per user, per month

For reference, consumer AI plans cost about \$20 a month

One set of environments serves all three, **and every new model needs new tasks**

Estimates, not signed prices. Lab prices are market rates from Epoch AI interviews with environment builders and AI labs (January 2026). Enterprise and personal prices are our assumptions.

Why us

Data vendors sell expert hours. **We build assets that compound with every customer**

	Expert-data vendors	SimReal
What compounds	Expert hours. Each project starts from scratch.	Proprietary environments, verifiers and maps of where models fail, plus the models we train in them. Every customer and loop adds to them.
Cost curve	Rises with headcount	Falls with scale: once a verifier exists, new data is generated and checked automatically
Proof of quality	Spot checks and reviewer judgment	Hidden tests, out-of-sample scoring, hash-committed (SHA-256) results and measured training gains (+12%)
Velocity	A new domain means new recruiting	7 environments in 14 days; \$3M ARR in 24 days
Network	Large expert pools, paid by the hour	Practitioners from top trading firms and 200K+ reachable experts, whose work our verifiers check

In short

They sell hours; we sell verified gains. **Their costs rise with scale; ours fall**

Network

Experts from the most competitive entry-level jobs and from research, **the work agents are now built to do**

Practitioners in our R&D have worked at top trading firms and markets, including



The most competitive entry-level jobs

Trading, finance, software and more

→ **Data for agents that streamline everyday work**

Research

AI, mathematics and quantitative research

→ **Data for agents that automate research**

200K+

Reachable experts

7,000+

Experts on the waitlist

21

Top universities

Students and alumni of



Logos show where practitioners have worked or studied; no endorsement implied.

Why Optiver PSI

A clear connection to market making, **with a long runway as AI moves into trading**

PSI looks for “a clear connection and long-term innovation potential for its core business” and invests people as well as capital.

01

Clear connection

Our flagship environment tests agents on market making, options, baskets and position limits, scored on risk-adjusted P&L on a held-out day.

02

Long-term potential

As AI agents move into research and trading, they must be tested out of sample before they touch capital. Environments and verifiers are that test.

03

People as well as capital

Optiver's market-structure experience would make our environments more realistic. Our CTO has quant experience at Optiver.

Raising a \$6M seed round

To scale environments, self-improvement research and delivery. business@simreal.co