



AI that improves itself in the real world


Real worlds rebuilt from real data, where AI labs, enterprises and personal agents train

This round

RMB 40M (paid in USD)

Business plan

September 2026



Overview

Give AI and personal agents a real world to fail, learn and **improve themselves**

What we do

Environments & AI data

For AI labs and enterprises: agent trajectories, expert and eval data; exclusive environments for RSI
Next: practice worlds for personal agents

Team

Quant founders, born 2005

Math at Cambridge, LSE and Duke
Jane Street, Citadel, Optiver, Millennium

Done so far

7 products in 14 days

Three worlds live: trading, AI research, event prediction; 5 public repos; no outside funding

Market

\$8.5B → \$700B

Labs buy \$8.5B a year of data and RL environments
The 2030 training market: ~\$700B, 82x today

Key result

+12%

Trained in Xitadel, Qwen3.8-27B traded up to 12% better on unseen trading days; replicated

This round

RMB 40M

Paid in USD; RMB 500M post-money
RMB 20M each for the revenue and RSI engines

Own an industry's best training world, own its best AI. **From trading to every industry: AI that improves itself.**

Team

From olympiads to Wall Street's toughest desks: **we know the gap between tests and real work**



Charles CEO

First employee at United Stables: helped take its U stablecoin from 0 to \$1.4B in a year; listed on Binance in a month

Ran institutional relations; worked on partnerships with SIG, DRW and others

Sole intern on HSBC's HKD stablecoin issuance; supported HKMA compliance end to end

X (Twitter) writer with millions of views

Hedge fund intern at Citadel, London

Born 2005 · LSE Mathematics · SCIE
USAMO qualifier

Henry CTO

Summer ML research at Cambridge with renowned statistician Po-Ling Loh (IMS Fellow)

Youngest undergraduate AI researcher at a Cambridge research center

Quant roles at Jane Street, Citadel, Optiver

Designed five benchmarks and RL environments

Born 2005 · Cambridge Math, First (Scholar) · SCIE
Top 30 worldwide, Cambridge math contest
UK Physics Olympiad Super Gold

Amaris COO

Data scientist at Millennium Hong Kong

The alt-data team's first-ever graduate hire

Helped launch Plug and Play's first Hong Kong event (with HK Science Park); hosted J&J

MedTech's tech summit (200+ attendees each)

Published author at 17; 100K+ engagements online

Born 2005 · Duke Math & Statistics · YK Pao School

Scale AI, Mercor and AfterQuery were founded by people aged about 20. Early investors made **17,000x, 40x, 1,800x**

Problem

AI can reason, **but it still can't do real work**

Real work has no answer key, only the world's reaction

	Surface check: passed		Real world: failed
Trading	The backtest only goes up	→	Loses money on unseen live markets
Software	All tests pass	→	Breaks in production
Finance	The books look finished	→	Month-end won't reconcile
Prediction	The analysis sounds convincing	→	Wrong once the result is in

To learn real work, AI needs **a world where real outcomes settle every move.**

Core insight

Every leap in AI came from a better feedback signal

Gen 1 · Internet data

Imitation

Models learn to copy

AI that can talk

Gen 2 · Human preference

Alignment

People rate the output

Useful assistants

Gen 3 · Answer keys

Verification

Checked against answers

Reasoning models

Gen 4 · Real-world feedback

Practice

AI acts in real situations; real outcomes settle every move

Autonomy and self-improvement

Just beginning

Acing tests isn't doing the job. AI's next leap comes from the real world's reaction.

Solution

AI does the task; **real outcomes become the next training signal**

- | | | |
|---------|-----------------|--|
| 01 | Environments | Real situations, rebuilt from real data |
| 02 | Outcome grading | Scored by real results, not another AI's opinion |
| 03 | Expert feedback | Professional judgment and grading standards |
| SimReal | | One environment for evals, data and training |

AI acts



The world responds with what actually happened



Settled by outcome; trajectory saved



Next round of training: self-improvement

Each round starts where the last one ended: that is **self-improvement (RSI)**.

Flagship proof · Xitadel

In two weeks we ran **the world's first self-improving market-making loop**

Real battlefield

Replays real trading days and order books, with data licensed through IMC Trading

Market as judge

The market settles every trade

Learns from results

Reviews each day's P&L and orders; each round beats the last

+12%

Open-source Qwen3.8-27B, after training, traded up to 12% better on unseen trading days

Controlled experiment, replicated across independent runs; logs available in diligence


Trading is the first world that works

Next: more strategies, more markets, more industries

Xitadel public preview

No model has passed the human baseline yet

Human baseline 80



Products

7 products shipped in 14 days, **5 of them in public repos**

Product	Domain	What it does	Status
Xitadel Xitadel-QuantBench	Trading	Replays real markets and settles on P&L; self-improvement proven	Open source · 102 stars
SimReal-MLBench Simreal-MLBench	AI research	60 research tasks across 7 data types, modeled on OpenAI's MLE-bench	Open source · 178 stars
FuturePredict Bench future-prediction-bench	Prediction	Predicts real events and trains on the outcome; live data	Partly open
MathmoBench MathmoBench	Math proofs	Makes AI prove its answers instead of guessing	Open source · 102 stars
Puzzle Benchmark hard-puzzle-benchmark	Logic	749 hand-written reasoning puzzles	Open source · 17 stars
Month-End Close	Finance	Trains AI to close the books with zero errors	Live · private
SWE-Forward	Software	Tests whether AI-written code survives the next release	Live · private

All built with **zero outside funding**.

400 stars in total · September 29, 2026



Data business

Environments and experts produce three kinds of AI data: **trajectories, expert data, evals**

01

Agent trajectories

A full record of AI doing a multi-step task: every action, the world's response, the result

Pass or fail set by real outcomes

For fine-tuning and RL

02

Expert data

From our network of students and alumni at 21 top universities: demonstrations, judgments and grading standards from real work

Entry-level roles to top research

For alignment, reward models and grading

03

Eval data

Private eval and held-out sets, used only for testing, never for training

Red-teamed before launch

For model acceptance and ongoing evals

Unlike human-only data vendors

Same source

Data, grading and training share one environment

Verifiable

Real outcomes decide pass or fail, not human opinion

Compounding

Where the model fails, the next batch of data goes

Labs spend \$8.5B a year on training data and RL environments: **we sell both.**

Next · Personal agents

Before an agent acts for someone, **it should practice somewhere real**

September 2026: the personal-agent race is on

Meta · Sep 8

Muse

Books, shops and schedules for you;
blocked by Amazon two weeks later

OpenAI · Sep 29

Dots

Always-on personal agents that
keep working in the background

Instinct · Sep 28

\$10B

Valuation: \$2.5B to \$10B in 33 days;
Sequoia, Benchmark, Coatue


Still falling short

16.1%

The best model completes 16.1% of real
freelance projects (Remote Labor Index)

What keeps the loop turning: **personal agents**

Dark: proven in Xitadel · Light: planned this round



Data rules Partners decide: pay to practice in private, or share de-identified runs for credits and a revenue share. Data reaches labs only with a license record.

Sandboxes keep agents out of trouble. **We teach them to get it right.**

Sources: Meta and OpenAI launches (Sep 8 and 29, 2026); Bloomberg (Sep 21, 2026); Instinct (Sep 28, 2026); Remote Labor Index (Scale AI and CAIS, Jul 2026). See A2.

Why now

Mercor grew revenue 27x and valuation 10x within 18 months: **the market is just starting**

Revenue growth

27x

Mercor gross run-rate: \$75M → \$2B in 16 months

Valuation growth

10x

Mercor valuation: \$2B → \$20B in 17 months (in talks)

Demand growth

14x

AI agent market forecast: \$7.9B → ~\$111B over 7 years

18x

Snorkel AI run-rate revenue in a year, now \$375M

14 months

AfterQuery: from founding to \$100M run-rate

5 months

AfterQuery valuation: \$300M → \$3.2B

\$29B

Scale AI valuation; Meta bought 49% for \$14.3B

Nearly half

of Scale AI's new training projects use RL environments

Before 2032

Public human text is projected to run out

Sources: Mercor (TechCrunch, Sacra, Dealroom, Bloomberg); AfterQuery (Business Wire, Forbes); Snorkel AI (company); Scale AI (Reuters, company blog); AI agents (Precedence Research); text stock (Epoch AI). See A2.

Market size

By 2030 the training market reaches **\$700B a year**, 82x today

Today · what labs buy

\$8.5B / year

Training data and RL environments, 50+ vendors

Top three by revenue (Surge: 2024; others run-rate)

Mercor		\$2.0B
Surge AI		\$1.2B
Snorkel AI		\$375M

2030 · AI economy

→ **~\$7T / year**

US (McKinsey)	\$2.9T
÷ US share of world GDP	25.7%
= world, same penetration	\$11.3T
= half penetration outside US	\$7.1T

→

2030 · Our market

\$700B / year

AI economy	\$7T
× training share	10%
= training market	\$700B

82x today

Training is the AI economy's R&D budget. Big tech spends 10–15% of revenue on R&D; we use 10%.

Today we sell to labs. By 2030, to **the whole AI economy**.

Sources: Menlo Ventures (Jul 2026); Mercor gross run-rate (Jun 2026), Surge revenue (2024), Snorkel run-rate (Sep 2026); McKinsey (Nov 2025); IMF (Apr 2026); company filings. Training share and non-US penetration are assumptions; see A2.

Competition

Others do one piece. We run **the whole loop that keeps AI improving**

Expert data platforms Mercor, Surge, AfterQuery	Expert demos and judgments, human-graded
China peers UniPat, Humanlaya	Mainly serve Chinese labs: expert data and evals
AI eval companies	Measure today's level; the score is the product
Labs in-house	Build only for work they already know
SimReal Environments, data, grading, RSI	One environment for data, grading and training Every result feeds the next round

China today Per Bloomberg, Alibaba led a \$300M round in UniPat at a \$2.5B valuation; Humanlaya closed a CDH-led Pre-A of several hundred million RMB; Alibaba, ByteDance and DeepSeek already buy from both.

Our edge

Own assets

Environment tooling, verifiers, a professional task library, accumulated failure cases

Stickiness

Every model upgrade needs new tasks, new data, new held-out sets and verifier upkeep

Speed

7 products in 14 days; self-improving market making in 2 weeks


Neutral

Owned by no lab: US and Chinese labs can both buy with confidence

Why us

The faster the frontier shifts, **the better for us**

Every wave was won by the fastest young team



Bet on us, and you bet on **every paradigm shift in AI.**

Business model

Every cycle brings a new delivery: **the deeper the partnership, the more we earn**

Three lines Environments · AI data (trajectories, experts, evals) · RSI service

Same model: AfterQuery hit \$100M run-rate in 14 months

Stage	What they buy	How we charge	Why they buy again
01 Paid pilot	Environment, data and eval pack for one capability	Fixed scope and acceptance criteria	Proves integration and training value
02 License	Environments, task sets, verifiers, agreed usage rights	Term, scope, exclusivity; data by volume	Wider domains and task coverage
03 Updates	New scenarios, tasks, trajectories, private held-out sets	Annual contract or batch orders	Old tasks saturate; new skills need new data
04 Joint RSI training	Training runs, ablations, transfer tests, hosted runs	Project or ongoing service contract	Closes the customer's next capability gap

Open source vs. paid

Open benchmarks: public, to build trust
Paid: private tasks, data, held-out sets, verifiers

Non-exclusive vs. exclusive

Non-exclusive: one environment, many labs
Exclusive: by domain and term, premium (Epoch AI)

Next · Personal agents

One world, two customers: labs and agents
Agent developers pay to practice; private by default

Traction

Products are public; **talks with frontier labs are under way**

2

Frontier labs in talks

7,000+

Experts on the waitlist

400

GitHub stars in total


September 29, 2026

5

Angels reached out to us

Frontier lab purchasing

Now



Traders who support our R&D



Optiver Δ

Polymarket

Logos show where supporters work; they do not imply endorsement by these firms.

Expert network

200,000+ reachable, verified experts drawn from students and alumni of 21 top universities: **our expert data source**

200K+ 21

Reachable, verified experts

Top universities

Selected universities in the network



Entry roles The work AI takes over first, in every industry

Top research The hardest problems and judgment calls

Moving up Members get promoted; the network climbs to senior levels



Each domain leaves assets behind: experts → data and standards → environments and verifiers → failure library → **reuse in the next domain**

Network coverage does not mean registration or delivery; university logos do not imply endorsement.

The round

Raising RMB 40M **at RMB 500M post-money**

Raising

RMB 40M

Paid in USD

Dilution

8%

Pre-money

RMB 460M

Post-money

RMB 500M

This round Environments: pilot → license → renewals; AI data at scale; RSI live in trading and prediction; worlds open to personal agents

Valuation references: top investors have already priced this space

China · UniPat

\$2.5B

Reported valuation (Sep 2026)

Alibaba led; Tencent, HongShan joined

China · Humanlaya

Hundreds of millions

Pre-A round size in RMB (Sep 2026)

CDH led; HongShan and others joined

Global · Applied Compute

\$100M → ~\$3B

Seed post → 14 months later (in talks)

\$20M seed round

Global · AfterQuery

\$300M → \$3.2B

Series A → 5 months later

Series A at \$100M run-rate

HongShan, CDH, Capital Today, BAI, Alibaba and Tencent are already in: **the category is validated.**

UniPat: reported valuation; Applied Compute ~\$3B: valuation in talks; Humanlaya: round size; others post-money. Sources in A2.

Use of funds

Two engines: revenue earns today's money; **RSI wins every industry**

01 Revenue engine · RMB 20M

Environments & data

Build once, license to many labs; data billed by volume

RMB 7M	Environments	Batch-build environments by industry
RMB 4M	Data	Agent trajectories, expert and eval data
RMB 6M	Delivery	Integration, acceptance, updates
RMB 3M	Sales & ops	Labs, enterprises, agent developers

Environment revenue = environments × licenses × price

Environment price \$20K–\$300K (Epoch AI)

3 months	6 months	12 months
First paid pilot	Licenses, renewals	Update contracts

Milestones are this round's targets.

02 RSI engine · RMB 20M

Self-improvement


The whole AI economy: in every industry, AI that improves itself

RMB 9M	Compute	Models train themselves, each round stronger
RMB 7M	Research	From trading to 10 industries
RMB 4M	Live trading	Real money and settlement; compliance

Trading RSI · Prediction RSI · 10 industries

A self-improving environment for every industry, in parallel

3 months	6 months	12 months
Trading RSI live	Prediction RSI live	Across 10 industries




Every industry's best AI will come from our worlds.

Three worlds live: trading, AI research, event prediction

SimReal · AI that improves itself in the real world

`business@simreal.co` · `simreal.co` · `x.com/simrealhq`



simreal

Confidential · For invited investors only · September 2026

Xitadel: what we tested and how

Environment	Replays real trading days and order books with data licensed through IMC Trading; the market settles every trade
Model	Open-source Qwen3.8-27B, before vs. after training
Test data	Real trading days the model had never seen
Result	Up to 12% better trading than the base model, replicated across independent runs (controlled experiment)
Public benchmark	Xitadel public preview: human baseline 80; best frontier model 77.28 (GPT 6); no model has passed the human baseline yet
Diligence	Run logs, metric definitions and scripts

Sources

Companies

Mercor run-rate: TechCrunch (Feb 2025, \$75M); CEO (Sep 2025, \$500M); Sacra (Dec 2025, \$760M); Mercor (early 2026, \$1B); Dealroom (Jun 2026, \$2B). All gross; experts take 60–70% (Bloomberg)

Mercor valuation: Series A \$250M (Sep 2024), Series B \$2B (Feb 2025), Series C \$10B (Oct 2025); a \$20B round in early talks (Bloomberg, Jul 9, 2026)

Surge AI: 2024 revenue \$1.2B (TechCrunch, Forbes)

Snorkel AI: Company and TechCrunch, Sep 22, 2026 (\$350M raised at a \$3.5B valuation; \$375M run-rate, up 18x in a year)

AfterQuery: Founded Feb 2025; Forbes, Sep 1, 2026 (18 months after YC, \$3.2B valuation); Business Wire (Apr 2026, Series A at \$300M, \$100M run-rate); YC profile (co-founder interned at Citadel Securities)

Scale AI: Founded in 2016 by Alexandr Wang, then 19 (Forbes); Meta bought 49% for \$14.3B, ~\$29B valuation (Reuters, Jun 2025); nearly half of new training projects involve RL environments (Scale AI blog, Feb 2026)

UniPat: Bloomberg, Sep 10, 2026 (Alibaba-led \$300M round, reported \$2.5B valuation, with Tencent and HongShan; terms may change; Alibaba, ByteDance, DeepSeek and others buy from UniPat and Humanlaya)

Humanlaya: Founded Oct 2025; Pre-A of several hundred million RMB on Sep 9, 2026, led by CDH with HongShan, Capital Today and BAI (Jiemian, Eastmoney)

Market and research

Training data spend today: Menlo Ventures (Jul 2026): 50+ vendors, ~\$8.5B combined

2030 AI economy: McKinsey, Agents, robots, and us (Nov 2025): US ~\$2.9T; IMF WEO (Apr 2026): US GDP \$32.4T, world \$126.3T; half the US penetration assumed elsewhere

Training market: 2030 AI economy × 10% (an assumption; big tech spends 10–15% of revenue on R&D)

AI agent market: Precedence Research (\$7.92B in 2025; 45.82% CAGR); 2028 and 2032 extrapolated

RL environment pricing: Epoch AI, An FAQ on RL environments (Jan 2026)

Public text stock: Epoch AI, Will we run out of data? (2024): exhausted 2026–2032 (80% CI)

Rounds and valuations: Applied Compute: Upstarts (Jun 2025, \$20M seed at \$100M post); The Information (Aug 2026, ~\$3B in talks)

Personal agents: Meta Muse launch (Sep 8, 2026; Axios, CNBC); Bloomberg (Sep 21, 2026): Amazon blocks Muse; OpenAI unveils always-on Dots at DevDay (Sep 29, 2026; TechCrunch, CNBC)

Instinct: Business Wire, Sep 28, 2026: \$1B Series C at a \$10B valuation from Sequoia, Benchmark and Coatue; valued at \$2.5B 33 days earlier

Remote Labor Index: Scale AI and CAIS, Jul 2026: the best model completes 16.1% of real freelance projects to a paying client's standard

Glossary

01

Training environment

A system where AI does real work again and again and learns from the results; an RL environment

02

Agent trajectory

The full record of AI doing a task: every action, the world's response, the result

03

Real feedback

Real outcomes settle each AI action: P&L, the books, whether code runs, whether an event happens

04

Self-improvement (RSI)

Recursive self-improvement: AI trains on its own results in real environments, stronger each round

05

Held-out set

Test-only tasks, never used in training, so models can't memorize answers

06

Red-teaming

Simulated cheating and attacks before launch, to find and fix grading loopholes

07

Controlled experiment

Change one thing (training in Xitadel or not) and compare before and after

08

Gross run-rate

Current revenue annualized, including payouts to experts