



AI that improves itself in the real world

Before a personal agent is trusted, it practices here

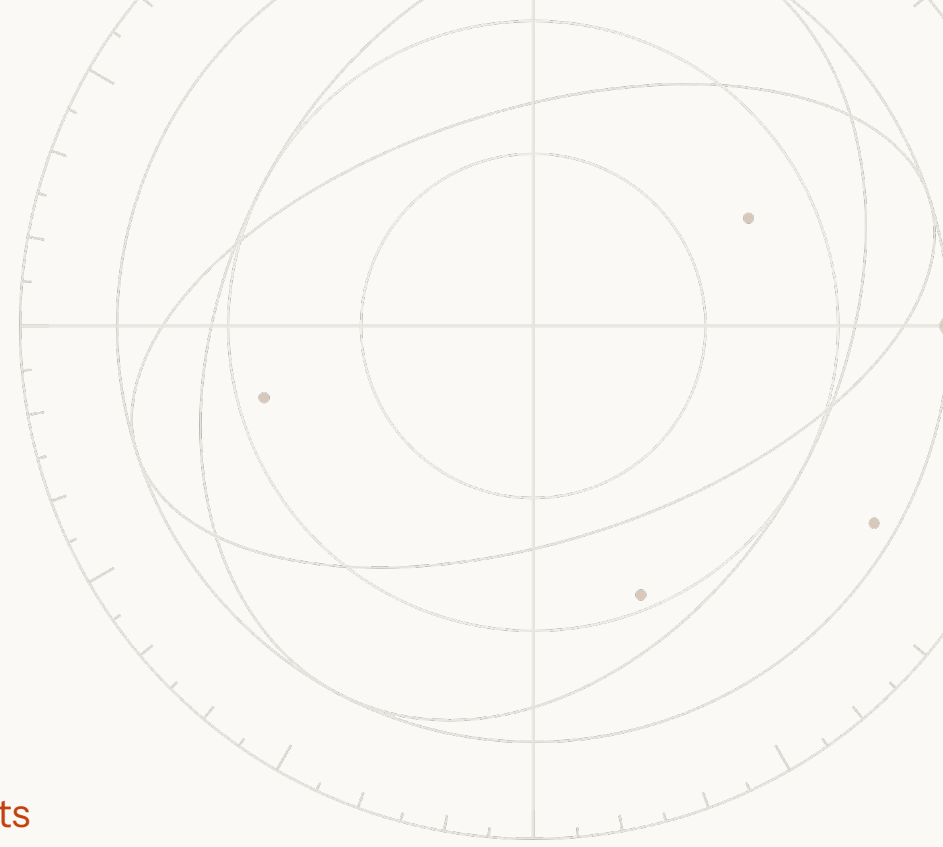
Starting with trading: training environments and data for AI labs and personal agents

This round

\$6M seed

Business plan

September 2026



Overview

Professional work like trading, turned into AI training environments: **scorable, reproducible, trainable**

What we do

Finance RL environments

Starting with trading: order books, simulator scoring
Sold to labs: licenses, agent trajectories, evals
Next: open practice to personal agents

Team

Quant founders

Math at Cambridge, LSE and Duke; born 2005
Jane Street, Citadel, Optiver, Millennium

Done so far

7 environments

Trading, prediction, AI research, math, reasoning,
finance, software · 5 public repos · no outside funding

Market

\$3.6M–\$42M a year

Finance alone, 17–24 buying labs
Bottom-up from frontier-lab pricing

Early evidence

+12%

Qwen3.8-27B, up to 12% better on unseen
competition trading days; best run, not a mean

This round

\$6M

At a \$75M post-money valuation

Win 2–3 paying labs in finance first, **then copy the method to the next domain.**

Team

From olympiads to Wall Street's toughest desks: **we know the gap between tests and real work**



Charles CEO

First employee at United Stables: helped take its U stablecoin from 0 to \$1.4B in a year; listed on Binance in a month

Ran institutional relations; worked on partnerships with SIG, DRW and others

Sole intern on HSBC's HKD stablecoin issuance; supported HKMA compliance end to end

X (Twitter) writer with millions of views

Hedge fund intern at Citadel, London

Born 2005 · LSE, Math · SCIE (Shenzhen)
USAMO qualifier

Henry CTO

Summer ML research at Cambridge with renowned statistician Po-Ling Loh (IMS Fellow)

Youngest undergraduate AI researcher at a Cambridge research center

Quant roles at Jane Street, Citadel, Optiver

Designed five benchmarks and RL environments

Born 2005 · Cambridge Math, First-Class (Scholar)
Top 30 worldwide, Cambridge math contest
UK Physics Olympiad Super Gold · SCIE (Shenzhen)

Amaris COO

Data scientist at Millennium Hong Kong

The alt-data team's first-ever graduate hire

Helped launch Plug and Play's first Hong Kong event (with HK Science Park); hosted J&J MedTech's tech summit (200+ attendees each)

Published author at 17; 100K+ engagements online

Born 2005 · Duke, Math & Statistics
YK Pao School (Shanghai)

Turned down return offers from top quant firms **to build this.**

Problem & solution

AI can reason but still can't do real work: **make real outcomes the next training signal**

Real work has no answer key, only the world's reaction

	Surface check: passed		Real world: failed
Trading	Backtest only goes up	→	Loses on unseen live markets
Software	All tests pass	→	Breaks in production
Finance	Books look finished	→	Month-end won't reconcile
Prediction	Analysis sounds right	→	Bets on it lose money

AI needs **a world where real outcomes settle every move**



01 Environments

Professional settings rebuilt by real rules

02 Scoring

Scored on real outcomes, not an AI's opinion

03 Expert input

Expert judgment and scoring standards

SimReal

One environment for evals, data and training

Each round keeps what real outcomes reward and retrains on it: **expert iteration today, RSI as the long-term goal.**

RSI: recursive self-improvement, where AI helps improve the next AI

Why now

Labs already pay for environments and expert data; **frontier models still lose money trading**

Labs are buying environments

\$1.5B+

Google's reported deal with RL-environment startup
Mechanize

Jul 2026

Mercor acquired RL-environment startup Deeptune

Expert data is scaling

27x

Mercor gross run-rate: \$75M → \$2B in 16 months

Tens of \$M

Paid by NVIDIA to Mercor in one quarter for expert data

18x

Snorkel AI run-rate growth in a year, to \$375M

Trading still loses

6/32

Alpha Arena: frontier models trading live; 6 of 32 runs
made money

2/7

PolyBench: 2 of 7 frontier models profitable on Polymarket

Sources: Mechanize (Business Insider); Mercor (TechCrunch, Dealroom, The Information, Forbes); Snorkel AI (Reuters); Alpha Arena (Nof1); PolyBench (arXiv).

Flagship proof · Xitadel

In two weeks we ran **the iterative training loop for market making**

One episode

Researches the order books for 6 or 12 hours with sandboxed tools, writes `Trader.run(state)`; no resubmits

Competition market

Replays IMC Prosperity order books tick by tick (matching rules on page 12)

Simulator as judge

Scored on P&L, drawdown and Sharpe; the formula is on page 12

Human benchmark

The best human competition strategy on the same held-out day scores 80

+12%


After training, the open model Qwen3.8-27B did up to 12% better on unseen competition days
Best run, not a mean

REPORTED WITHIN 3 MONTHS

Each task × 10 seeds, paired before/after: mean, 95% CI, paired p-value, per-task change; reserve tasks as a second held-out set; event prediction checked with Brier scores

Xitadel public preview · 7 tasks

GPT 6 matches humans on 5 of 7 tasks; multi-asset only 51



task_04: GLM 5.3 96.03, GPT 6 89.53; Kimi K3 and DeepSeek V4 Pro 0
Overall = mean of 7 tasks; public runs used 2/4-hour pilot limits (standard 6/12)

Limits: the competition market is made of trading bots, not exchange data; replay ignores the price impact of our orders; real market data comes next. Data: IMC Prosperity 3 and 4 open-source backtester resources (MIT).

Only numbers that reproduce: **raw logs, metric definitions, scripts and result hashes in diligence.**

Source: Xitadel-QuantBench public repo (README, REPORT, SCORING).



Products & data

7 benchmarks and environments, from trading to software; **5 open-sourced**

Product	Domain	What it does	Status
Xitadel Xitadel-QuantBench	Trading	Flagship market-making environment (page 6)	Open · 102 stars
SimReal-MLBench Simreal-MLBench	AI research	60 research tasks, 7 data types; modeled on OpenAI's MLE-bench	Open · 178 stars
FuturePredict Bench future-prediction-bench	Prediction	Predicts real events, trains on outcomes once resolved; live data	Partly open
MathmoBench MathmoBench	Math proofs	Makes AI prove answers instead of guessing them	Open · 102 stars
Puzzle Benchmark hard-puzzle-benchmark	Reasoning	749 hand-written reasoning puzzles	Open · 17 stars
Month-End Close	Finance close	Has AI close the books with zero errors	Live · private
SWE-Forward	Software	Tests whether AI-written code survives the next version	Live · private

Every environment yields three kinds of data

Agent trajectories

Every action, the environment's feedback and the outcome, labeled by real results
For fine-tuning and RL

Expert data

Demonstrations, judgments and scoring standards from real work, via our expert network (page 8)
For alignment, reward models and grading

Self-extending: where the model fails, the next batch fills in

Eval data

Private eval and held-out sets: for testing only, never for training; red-teamed
For model acceptance and ongoing evals

Stars as of September 29, 2026



Expert network

7,000+ experts on the waitlist, students and alumni of 21 top universities: **our expert data source**

7,000+ 21

On the waitlist (200K+ reachable) Top universities

Selected universities in the network

Entry roles The work AI takes over first, in every industry

Top research The hardest problems and judgment calls

Moving up Members get promoted; the network climbs to senior levels



UNIVERSITY OF
CAMBRIDGE



Yale University



THE UNIVERSITY OF
CHICAGO



Each domain leaves assets behind: experts → data and standards → environments and verifiers → failure library → **reuse in the next domain**

Network coverage does not mean registration or delivery; university logos do not imply endorsement.

Next · Personal agents

Before an agent acts for someone, **it should practice somewhere real**

Sep 2026: personal agents start spending for people (Instinct’s founder cites \$1B+ a year, about half travel) and making mistakes for them

Muse

Meta, launched Sep 8;
No. 1 on the US App Store in 10 days

Dots

OpenAI, unveiled Sep 29;
always on in the cloud, working for you

\$10B

Instinct’s valuation, 4x in 33 days;
Sequoia, Benchmark, Coatue

\$200

Cancellation fee a user paid after
Instinct booked without asking

How we serve them: the same environments, a second customer

What opens

Trading and event prediction; held-out sets stay closed to practice

Pricing

Per practice round, \$2–\$10 by task difficulty and length; private by default

Data rules

De-identified trajectories can be shared for credits and revenue share; data sent to labs carries consent records

Milestone

First 10 agent teams onboarded within 6 months

Sandboxes keep agents out of trouble. **We teach them to get it right.**

Sources: Meta and OpenAI launches (Sep 2026); TechCrunch (Sep 25, 2026); Instinct (founder podcast, Sep 2026; Reuters, Sep 28, 2026; The Atlantic, CNN, Sep 2026).

Market size

Bottom-up from frontier-lab pricing: finance alone is **about \$3.6M–\$42M a year**

① Epoch prices × one lab’s volume

\$200–\$2,000

Per task; environments \$20K–\$300K each, finance priced at the upper-middle

Finance envs: 2–5 × \$100K–\$300K	\$200K–\$1.5M
Tasks: 500–1,500 × \$500–\$2,000	\$250K–\$3M
= Per US frontier lab, per year	\$450K–\$4.5M

② Buyers × annual spend per lab

→ 17–24 labs

Labs doing agent post-training, in three budget tiers

US frontier: 4–7 × \$450K–\$4.5M	\$1.8M–\$31.5M
New labs: 5–7 × \$200K–\$800K	\$1M–\$5.6M
China: 8–10 × \$100K–\$500K	\$800K–\$5M
≈ Finance alone, per year	\$3.6M–\$42M

→

③ If 2 more domains match finance (est.)

\$10.8M–\$126M

Finance alone, per year \$3.6M–\$42M

Trading + prediction + AI research × 3

= Three domains, per year \$10.8M–\$126M

Context: training-data and RL-environment vendors earn about \$8.5B a year
Source: Deedy Das, Jul 2026

One domain is worth millions to tens of millions a year; **each new scorable domain adds another.**

Sources: task and environment prices and contract sizes (six to seven figures a quarter): Epoch AI (Jan 2026, 18 interviews), SemiAnalysis (Jan 2026); vendor revenue: Deedy Das (Jul 2026), context only.
Estimates: finance priced at the upper-middle (task floor ~\$500, per Mechanize’s ~\$2,400 of training compute per task); volume per lab (China benchmarked to UniPat orders), tier counts and the three-domain extrapolation are our assumptions.
Buyers: OpenAI, Google DeepMind, Meta, xAI and other US frontier labs; Thinking Machines, Reflection AI and other new labs; Alibaba, ByteDance, DeepSeek, Moonshot and other Chinese labs.

Competition

Others build one piece; **we build the full loop that keeps AI improving**

Expert data platforms Mercor, Surge, Handshake, AfterQuery	Expert demos and judgment, scored by human opinion Moving into RL environments
Peers in China UniPat, Humanlaya	Expert data, eval environments and benchmarks; mainly domestic labs
AI trading evals Nof1 (Alpha Arena)	Live trading leaderboard: evaluates, doesn't train
Labs in-house	Only the domains they know
SimReal Environments, data, scoring, training	Every result feeds the next round; AI keeps improving

In China Bloomberg reports Alibaba plans to lead a \$300M round in UniPat at a \$2.5B valuation; Humanlaya closed a pre-A of several hundred million RMB led by CDH; Alibaba, ByteDance, DeepSeek and others have bought data or services from both.

Our edge

Finance know-how

A quant team designs the matching and the reward: trading discipline written into the scorer (next page)

Verifiable scoring

Scored by the simulator, not by opinion; human benchmark from 6 competition teams' public submissions

Evaluate and train

One environment for evals and training: Qwen3.8-27B up to 12% better after training (page 6)

Own assets

Matching and scoring on our own engine; a 7,000+ expert waitlist (page 8)

Why us

Labs buy finance know-how by the hour; **we built it into the environment and the score**

OpenAI pays ex-bankers \$150 an hour; xAI hires quant-trading experts (Bloomberg, xAI job posts)

Realistic matching

Naive backtests assume resting orders fill, so models learn the backtest's loopholes. We queue orders, split crossing and resting fills, and attribute each fill to the book.

Risk in the reward

$A = \max(0, P\&L - 0.1 \times DD) \times \text{Sharpe factor}$
After fixing 3 flaws in a human strategy, P&L rose 63% but Sharpe fell 7.27 $\rightarrow$ 3.03: return and risk diverge; the reward counts both

No gaming

The last full day is hidden from the agent; the full report's hash is published up front; deterministic: same strategy, same score; private reserve tasks replace exposed days

Quant experience

All three founders have quant-firm experience (page 3); the team designed the three rules above

What buyers want

Interviewees repeatedly name it a key quality test: no high score by gaming
Epoch AI, 18 interviews

Speed

7 environments in 14 days, no outside funding

Spotting the loophole in a backtest **is basic quant discipline.**

Business model & traction

Products are public; revenue starts with paid pilots: **the deeper the work, the bigger the contract**

2

frontier labs in talks

Three lines Training environments · AI data (agent, expert, eval) · Iterative training

Same model: AfterQuery reached a \$100M run-rate in 14 months

Stage	What the client buys	Length	Price	Why they keep buying
Free eval	Public preview: 7 tasks, training data, SDK, scoring rules	1-2 weeks	Free	Scores win clients; benchmarked to the best human
Paid pilot	Custom tasks, held-out days, reward API; data with a license chain	8-12 weeks	\$50K-\$150K	Before/after training results on agreed metrics
Annual license	Full environments, matching engine, counterparty models, reserve tasks, expert data	12 months	\$400K-\$2M	Tasks and held-out sets refresh quarterly; old tasks saturate
Joint training	Training runs, ablations, transfer tests, managed runs	Project or ongoing service contract		Solves the client's next capability gap

Unit economics (est.): a public environment takes about 6 person-days; from the second lab on, it is reused as is. Buyer concentration: extend to trading firms.

Supported by practitioners at trading firms



Logos show where supporters work; they do not imply endorsement by the firms.

The round & use of funds

Two engines: revenue earns today's money; **R&D makes the training gains solid**

Raising (seed)

\$6M

Dilution

8%

Pre-money

\$69M

Post-money

\$75M

01 Revenue engine · \$3M

Environments & data

Build once, license to many labs; data billed by volume

\$1.05M	Environments	\$600K	Data
\$900K	Delivery	\$450K	Sales & ops

5 environment engineers · 3 delivery · 2 sales; 24-month runway

3 months	6 months	12 months
First paid pilot	Second license	2–3 paying labs \$1M–\$3M ARR

02 R&D engine · \$3M

Iterative training

Trading → prediction → finance close: prove the gain, then scale

\$1.35M	Compute	\$1.05M	Research team
\$600K	Live trading, compliance		

3 researchers · ~600K GPU hours

3 months	6 months	12 months
Significance test, trading	Gains in event prediction	Small live-capital test

Valuation references: seed-stage peers at \$100M–\$1B post-money

Applied Compute \$100M

Seed post-money · Jun 2025

Mirendil \$1B

Seed post-money · Jun 2026


Nof1 \$15M

Raised · May 2026

Chaoyan 超衍智能 ~\$59M

Angel rounds raised · Sep 2026

Applied Compute, Mirendil: post-money; Nof1, Chaoyan: amount raised. Round: RMB 40M at RMB 500M post-money, at 6.7351 RMB/USD (Sep 30, 2026 central parity). Data stored and delivered by client region; Cayman–Hong Kong–onshore structure before closing.




Every industry's best AI will come from our worlds.

Three worlds live: trading, AI research, event prediction

SimReal · AI that improves itself in the real world

`business@simreal.co` · `simreal.co` · `x.com/simrealhq`



simreal

Confidential · For invited investors only · September 2026