



# AI that improves itself in the real world

Before a personal agent is trusted, it practices here

Starting with trading: training environments and data for AI labs and personal agents


---

This round

\$6M seed

Business plan

September 2026



# Overview

Professional work like trading, turned into AI training environments: **scorable, reproducible, trainable**

What we do

## Finance RL environments

Professional work like trading, made scorable  
Sold to labs: licenses, agent trajectories, evals  
Next: open practice to personal agents

Team

## Quant founders

Math at Cambridge, LSE and Duke; born 2005  
Jane Street, Citadel, Optiver, Millennium

Done so far

## 7 environments

Trading, prediction, AI research, math, reasoning,  
finance, software · 5 public repos · no outside funding

Market

**\$3.6M–\$42M a year**

Finance alone, 17–24 buying labs  
Bottom-up from frontier-lab pricing

Early evidence

**+12%**

Qwen3.8-27B, up to 12% better on unseen  
competition trading days; best run, not a mean

This round

**\$6M**

At a \$75M post-money valuation

Win 2–3 paying labs in finance first, **then copy the method to the next domain.**

# Team

From olympiads to Wall Street's toughest desks: **we know the gap between tests and real work**



## Charles CEO

First employee at United Stables: helped take its U stablecoin from 0 to \$1.4B in a year; listed on Binance in a month

Ran institutional relations; worked on partnerships with SIG, DRW and others

Sole intern on HSBC's HKD stablecoin issuance; supported HKMA compliance end to end

X (Twitter) writer with millions of views

Hedge fund intern at Citadel, London

---

Born 2005 · LSE, Math · SCIE (Shenzhen)  
USAMO qualifier

## Henry CTO

Summer ML research at Cambridge with renowned statistician Po-Ling Loh (IMS Fellow)

Youngest undergraduate AI researcher at a Cambridge research center

Quant roles at Jane Street, Citadel, Optiver

Designed five benchmarks and RL environments

---

Born 2005 · Cambridge Math, First-Class (Scholar)  
Top 30 worldwide, Cambridge math contest  
UK Physics Olympiad Super Gold · SCIE (Shenzhen)

## Amaris COO

Data scientist at Millennium Hong Kong

The alt-data team's first-ever graduate hire

Helped launch Plug and Play's first Hong Kong event (with HK Science Park); hosted J&J MedTech's tech summit (200+ attendees each)

Published author at 17; 100K+ engagements online

---

Born 2005 · Duke, Math & Statistics  
YK Pao School (Shanghai)

Turned down return offers from top quant firms **to build this.**

# Problem & solution

AI can reason but still can't do real work: **make real outcomes the next training signal**

Real work has no answer key, only the world's reaction

|            | Surface check: passed |   | Real world: failed           |
|------------|-----------------------|---|------------------------------|
| Trading    | Backtest only goes up | → | Loses on unseen live markets |
| Software   | All tests pass        | → | Breaks in production         |
| Finance    | Books look finished   | → | Month-end won't reconcile    |
| Prediction | Analysis sounds right | → | Bets on it lose money        |

AI needs **a world where real outcomes settle every move**



## 01 Environments

Professional settings rebuilt by real rules

## 02 Scoring

Scored on real outcomes, not an AI's opinion

## 03 Expert input

Expert judgment and scoring standards

## SimReal

One environment for evals, data and training

Each round keeps what real outcomes reward and retrains on it: **expert iteration today, RSI as the long-term goal.**

RSI: recursive self-improvement, where AI helps improve the next AI and each round's real outcomes become the next round's training data

# Why now

Labs already pay for environments and expert data; **frontier models still lose money trading**

---

Labs are buying environments

**\$1.5B+**

Google's reported deal with RL-environment startup  
Mechanize

---

**Jul 2026**

Mercor acquired RL-environment startup Deeptune

---

Expert data is scaling

**27x**

Mercor gross run-rate: \$75M → \$2B in 16 months

---

**Tens of \$M**

Paid by NVIDIA to Mercor in one quarter for expert data

---

**18x**

Snorkel AI run-rate growth in a year, to \$375M

---

Trading still loses

**6/32**

Alpha Arena: frontier models trading live; 6 of 32 runs  
made money

---

**2/7**

PolyBench: 2 of 7 frontier models profitable on Polymarket

Sources: Mechanize (Business Insider); Mercor (TechCrunch, Dealroom, The Information, Forbes); Snorkel AI (Reuters); Alpha Arena (Nof1); PolyBench (arXiv).

# Flagship proof · Xitadel

In two weeks we ran **the iterative training loop for market making**

|                           |                                                                                                                                                                 |
|---------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>One episode</b>        | Gets historical order books and product rules, researches for 6 or 12 hours with sandboxed tools, writes <code>Trader.run(state)</code> ; submissions are final |
| <b>Competition market</b> | Replays IMC Prosperity order books tick by tick; each fill is attributed to the book at order time                                                              |
| <b>Simulator as judge</b> | Scored on P&L, drawdown and Sharpe; the formula is on page 12                                                                                                   |
| <b>Human benchmark</b>    | The best human competition strategy on the same held-out day scores 80                                                                                          |

## +12%

After training, the open model Qwen3.8-27B did up to 12% better on unseen competition days  
**Best run, not a mean**

### REPORTED WITHIN 3 MONTHS

Each task × 10 random seeds, paired before/after: mean, 95% CI, paired p-value, per-task change; reserve tasks as a second held-out set; FuturePredict settles on real events, with before/after Brier scores


Limits: the competition market is made of trading bots, not exchange data; replay ignores the price impact of our orders; real market data comes next. Data: IMC Prosperity 3 and 4 open-source backtester resources (MIT).

Only numbers that reproduce: **raw logs, metric definitions, scripts and result hashes in diligence.**

Source: Xitadel-QuantBench public repo (README, REPORT, SCORING).

Xitadel public preview · 7 tasks

**GPT 6 matches humans on 5 of 7 tasks; multi-asset only 51**



task\_04: GLM 5.3 96.03, GPT 6 89.53; Kimi K3 and DeepSeek V4 Pro 0  
Overall = mean of 7 tasks; public runs used 2/4-hour pilot limits (standard 6/12)

# Products & data

7 benchmarks and environments in 14 days, 5 public repos; **no outside funding**

| Product                                               | Domain        | What it does                                                                     | Status                  |
|-------------------------------------------------------|---------------|----------------------------------------------------------------------------------|-------------------------|
| <b>Xitadel</b><br>Xitadel-QuantBench                  | Trading       | Replays competition order books; scores P&L, drawdown, Sharpe; training loop run | <b>Open</b> · 102 stars |
| <b>SimReal-MLBench</b><br>Simreal-MLBench             | AI research   | 60 research tasks, 7 data types; modeled on OpenAI's MLE-bench                   | <b>Open</b> · 178 stars |
| <b>FuturePredict Bench</b><br>future-prediction-bench | Prediction    | Predicts real events, trains on outcomes once resolved; live data                | <b>Partly open</b>      |
| <b>MathmoBench</b><br>MathmoBench                     | Math proofs   | Makes AI prove answers instead of guessing them                                  | <b>Open</b> · 102 stars |
| <b>Puzzle Benchmark</b><br>hard-puzzle-benchmark      | Reasoning     | 749 hand-written reasoning puzzles                                               | <b>Open</b> · 17 stars  |
| <b>Month-End Close</b>                                | Finance close | Has AI close the books with zero errors                                          | Live · private          |
| <b>SWE-Forward</b>                                    | Software      | Tests whether AI-written code survives the next version                          | Live · private          |

Every environment yields three kinds of data

## Agent trajectories

Every action, the environment's feedback and the outcome, labeled by real results  
For fine-tuning and RL

## Expert data

Demonstrations, judgments and scoring standards from real work, entry roles to top research  
For alignment, reward models and grading

**Self-extending: where the model fails, the next batch fills in**

## Eval data

Private eval and held-out sets, never used in training; red-teamed before release  
For model acceptance and ongoing evals

Stars as of September 29, 2026

# Expert network

7,000+ experts on the waitlist, students and alumni of 21 top universities: **our expert data source**

7,000+ 21

On the waitlist (200K+ reachable) Top universities

Entry roles The work AI takes over first, in every industry

Top research The hardest problems and judgment calls

**Moving up** Members get promoted; the network climbs to senior levels

Selected universities in the network



Each domain leaves assets behind: experts → data and standards → environments and verifiers → failure library → **reuse in the next domain**

Network coverage does not mean registration or delivery; university logos do not imply endorsement.

# Next · Personal agents

Before an agent acts for someone, **it should practice somewhere real**

Sep 2026: personal agents start spending for people (Instinct’s founder cites \$1B+ a year, about half travel) and making mistakes for them

Muse

Meta, launched Sep 8;  
No. 1 on the US App Store in 10 days

Dots

OpenAI, unveiled Sep 29;  
always on in the cloud, working for you

\$10B

Instinct’s valuation, 4x in 33 days;  
Sequoia, Benchmark, Coatue

\$200

Cancellation fee a user paid after  
Instinct booked without asking

How we serve them: the same environments, a second customer

**What opens**

Trading and event prediction; held-out sets stay closed to practice

**Pricing**

Per practice round, \$2–\$10 by task difficulty and length; private by default

**Data rules**

De-identified trajectories can be shared for credits and revenue share; data sent to labs carries consent records

**Milestone**

First 10 agent teams onboarded within 6 months

Sandboxes keep agents out of trouble. **We teach them to get it right.**

Sources: Meta and OpenAI launches (Sep 2026); TechCrunch (Sep 25, 2026); Instinct (founder podcast, Sep 2026; Reuters, Sep 28, 2026; The Atlantic, CNN, Sep 2026).

# Market size

Bottom-up from frontier-lab pricing: finance alone is **about \$3.6M–\$42M a year**

① Epoch prices × one lab's volume

**\$200–\$2,000**

Per task; environments \$20K–\$300K each, finance priced at the upper-middle

Finance envs: 2–5 × \$100K–\$300K **\$200K–\$1.5M**

Tasks: 500–1,500 × \$500–\$2,000 **\$250K–\$3M**

**= Per US frontier lab, per year \$450K–\$4.5M**

② Buyers × annual spend per lab

→ **17–24 labs**

Labs doing agent post-training, in three budget tiers

US frontier: 4–7 × \$450K–\$4.5M **\$1.8M–\$31.5M**

New labs: 5–7 × \$200K–\$800K **\$1M–\$5.6M**

China: 8–10 × \$100K–\$500K **\$800K–\$5M**

**≈ Finance alone, per year \$3.6M–\$42M**

→

③ Three live domains · per year (extrapolated)

**\$10.8M–\$126M**

Finance alone, per year **\$3.6M–\$42M**

× Event prediction, AI research × 3

**= Three domains, per year \$10.8M–\$126M**

Context: training-data and RL-environment vendors earn about \$8.5B a year (Deedy Das, Jul 2026)

**Finance is only the first domain**

One domain is worth millions to tens of millions a year; **each new scorable domain multiplies it.**

Sources: task and environment prices and contract sizes (six to seven figures a quarter): Epoch AI (Jan 2026, 18 interviews), SemiAnalysis (Jan 2026); vendor revenue: Deedy Das (Jul 2026), context only.

Estimates: finance priced at the upper-middle (task floor ~\$500, per Mechanize's ~\$2,400 of training compute per task); volume per lab, tier counts (China benchmarked to UniPat orders) and the three-domain extrapolation are our assumptions. Buyers: OpenAI, Google DeepMind, Meta, xAI and other US frontier labs; Thinking Machines, Reflection AI and other new labs; Alibaba, ByteDance, DeepSeek, Moonshot and other Chinese labs.

# Competition

Others build one piece; **we build the full loop that keeps AI improving**

|                                                                      |                                                                                   |
|----------------------------------------------------------------------|-----------------------------------------------------------------------------------|
| <b>Expert data platforms</b><br>Mercor, Surge, Handshake, AfterQuery | Expert demos and judgment, scored by human opinion<br>Moving into RL environments |
| <b>Peers in China</b><br>UniPat, Humanlaya                           | Expert data, eval environments and benchmarks; mainly domestic labs               |
| <b>AI trading evals</b><br>Nof1 (Alpha Arena)                        | Live trading leaderboard: evaluates, doesn't train                                |
| <b>Labs in-house</b>                                                 | Only the domains they know                                                        |
| <b>SimReal</b><br><b>Environments, data, scoring, training</b>       | Every result feeds the next round; AI keeps improving                             |

**In China** Bloomberg reports Alibaba plans to lead a \$300M round in UniPat at a \$2.5B valuation; Humanlaya closed a pre-A of several hundred million RMB led by CDH; Alibaba, ByteDance, DeepSeek and others have bought data or services from both.

## Our edge

### Finance know-how

A quant team designs the matching and the reward: trading discipline written into the scorer (next page)

### Verifiable scoring

Scored by the simulator, not by opinion; human benchmark from 6 competition teams' public submissions

### Evaluate and train

One environment for evals and training: eval data, agent trajectories and iterative training, delivered together

### Own assets

Matching and scoring on our own engine; reserve tasks and unpublished held-out days; 7 environments live

# Why us

Labs buy finance know-how by the hour; **we built it into the environment and the score**

OpenAI pays ex-bankers \$150 an hour; xAI hires quant-trading experts (Bloomberg, xAI job posts)

## Realistic matching

Naive backtests assume resting orders fill,  
so models learn the backtest's loopholes.  
We queue orders, split crossing and resting  
fills, and attribute each fill to the book.

## Risk in the reward

Score =  $\max(0, \text{P\&L} - 0.1 \times \text{DD}) \times \text{Sharpe factor}$   
After fixing 3 flaws in a human strategy,  
P&L rose 63% but Sharpe fell 7.27  $\rightarrow$  3.03:  
return and risk diverge, so risk is priced in

## No gaming

The last full day is hidden from the agent;  
the full report's hash is published up front;  
deterministic: same strategy, same score;  
private reserve tasks replace exposed days

## Quant experience

Jane Street, Citadel, Optiver (quant);  
Millennium data science; Cambridge, LSE, Duke math

## What buyers want

Labs choosing environments keep citing one key test: no  
high score by gaming (Epoch AI, 18 interviews)

## Speed

7 benchmarks and environments in 14 days, no outside  
funding

Spotting the loophole in a backtest **is basic quant discipline.**

# Business model & traction

Products are public; revenue starts with paid pilots: **the deeper the work, the bigger the contract**

2

frontier labs in talks

Three lines   Training environments · AI data (agent, expert, eval) · Iterative training

Same model: AfterQuery reached a \$100M run-rate in 14 months

| Stage          | What the client buys                                                                              | Length                              | Price                                   | Why they keep buying                                          |
|----------------|---------------------------------------------------------------------------------------------------|-------------------------------------|-----------------------------------------|---------------------------------------------------------------|
| Free eval      | Public preview: 7 tasks, training data, SDK, scoring rules                                        | 1-2 weeks                           | Free                                    | Scores win clients; benchmarked to the best human             |
| Paid pilot     | Custom tasks, held-out days, reward API; own or licensed data with a license chain                | 8-12 weeks                          | \$50K-\$150K                            | Before/after training results on agreed metrics               |
| Annual license | Full environments, matching engine, counterparty models, reserve tasks, trajectories, expert data | 12 months                           | \$400K-\$2M                             | Tasks and held-out sets refresh quarterly; old tasks saturate |
| Joint training | Training runs, ablations, transfer tests, managed runs                                            | Project or ongoing service contract | Solves the client's next capability gap |                                                               |

Unit economics (est.): a public environment takes about 6 person-days; a second lab reuses the same environment. Buyer concentration: license each environment to many labs and to trading firms.

Supported by practitioners at trading firms



Logos show where supporters work; they do not imply endorsement by the firms.

# The round & use of funds

Two engines: revenue earns today's money; **R&D makes the training gains solid**

Raising (seed)

**\$6M**

Dilution

**8%**

Pre-money

**\$69M**

Post-money

**\$75M**

**01 Revenue engine · \$3M**

## Environments & data

Build once, license to many labs; data billed by volume

|                |              |               |             |
|----------------|--------------|---------------|-------------|
| <b>\$1.05M</b> | Environments | <b>\$600K</b> | Data        |
| <b>\$900K</b>  | Delivery     | <b>\$450K</b> | Sales & ops |

5 environment engineers · 3 delivery · 2 sales; 24-month runway

|                  |                 |                                |
|------------------|-----------------|--------------------------------|
| <b>3 months</b>  | <b>6 months</b> | <b>12 months</b>               |
| First paid pilot | Second license  | 2–3 paying labs, \$1M–\$3M ARR |

**02 R&D engine · \$3M**

## Iterative training

Trading → prediction → finance close: prove the gain, then scale

|                |                          |                |               |
|----------------|--------------------------|----------------|---------------|
| <b>\$1.35M</b> | Compute                  | <b>\$1.05M</b> | Research team |
| <b>\$600K</b>  | Live trading, compliance |                |               |

3 researchers · ~600K GPU hours

|                                  |                           |                         |
|----------------------------------|---------------------------|-------------------------|
| <b>3 months</b>                  | <b>6 months</b>           | <b>12 months</b>        |
| Trading gains: significance test | Gains in event prediction | Small live-capital test |

Valuation references: seed-stage peers at \$100M–\$1B post-money

**Applied Compute \$100M**

Seed post-money · Jun 2025

**Mirendil \$1B**

Seed post-money · Jun 2026


**Nof1 \$15M**

Raised · May 2026

**Chaoyan 超衍智能 ~\$59M**

Angel rounds raised · Sep 2026

Applied Compute, Mirendil: post-money; Nof1, Chaoyan: amount raised. Round: RMB 40M at RMB 500M post-money, at 6.7351 RMB/USD (Sep 30, 2026 central parity). Data stored and delivered by client region; Cayman–Hong Kong–onshore structure before closing.




# Every industry's best AI will come from our worlds.

Three worlds live: trading, AI research, event prediction

**SimReal** · AI that improves itself in the real world

---

`business@simreal.co` · `simreal.co` · `x.com/simrealhq`



simreal

Confidential · For invited investors only · September 2026