



AI that improves itself in the real world

Before a personal agent is trusted, it practices here

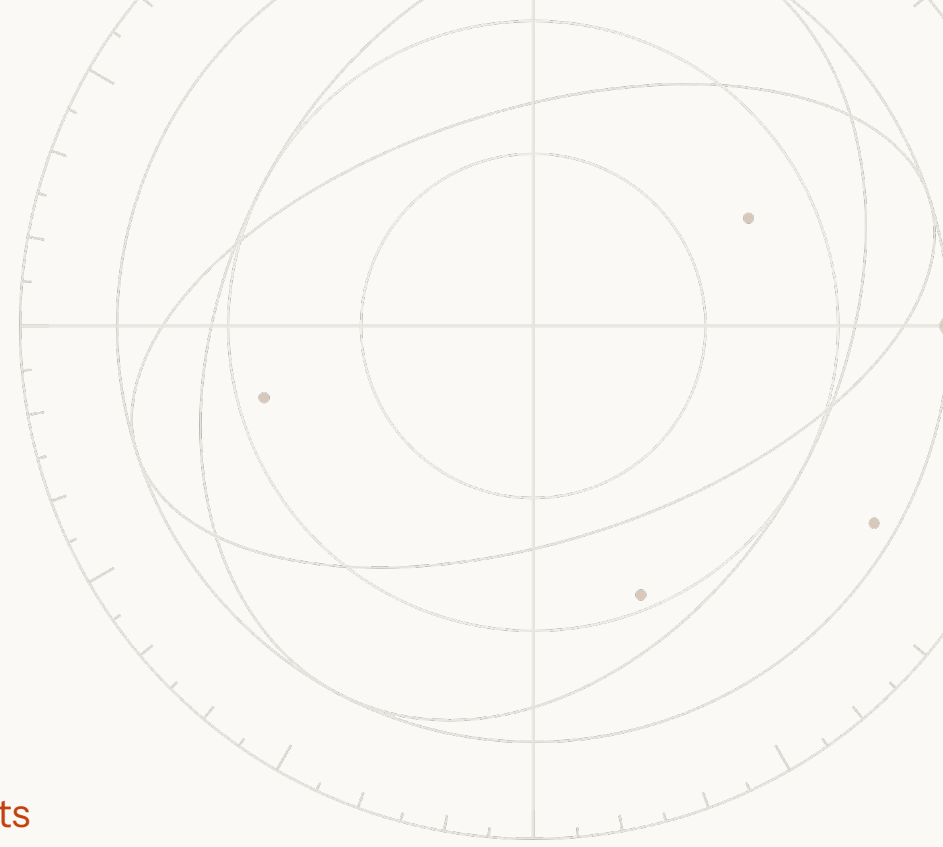
Starting with trading: training environments and data for AI labs and personal agents

This round

\$6M seed

Business plan

September 2026



Overview

SimReal builds training grounds for AI labs: **live, with results, with buyers in talks**

7 environments

Live in 14 days; 5 open-sourced; no outside funding

+12%

Best gain of an open model trained in our trading environment

2 labs

Frontier labs in purchase talks

7,000+

Experts signed up, from 21 top universities

\$3.6M–\$42M

Finance alone, per year, built up from lab prices

\$6M

This round, at a \$75M post-money valuation

Win 2–3 paying labs in finance, **then copy the method to the next domain.**

Team

From olympiads to Wall Street desks: **we know the gap between tests and real work**



Charles CEO

First employee at United Stables: U stablecoin from 0 to \$1.4B in a year

Sole intern on HSBC's HKD stablecoin; supported HKMA compliance

Hedge fund intern at Citadel, London

LSE Math · USAMO qualifier · born 2005

Henry CTO

Quant roles at Jane Street, Citadel, Optiver

Cambridge ML research with Po-Ling Loh (IMS Fellow)

Designed five benchmarks and RL environments

Cambridge Math, First-Class (Scholar) · born 2005

Amaris COO

Data scientist, Millennium Hong Kong

First-ever graduate hire on the alt-data team

Published author at 17; 100K+ engagements

Duke Math & Statistics · born 2005

Turned down return offers from top quant firms **to build this.**

Problem

AI can reason, **but it still can't do real work**

Passes the test, fails the job

Trading	Backtest only goes up	→ Loses money live
Software	All tests pass	→ Breaks in production
Finance	Books look done	→ Month-end won't close

6/32

Frontier models trading real money: 6 of 32 runs made a profit

→ **Trading is unsolved for AI, and right or wrong is easy to check**

51

GPT 6 on our multi-asset trading task (best human = 80)

→ **Our environment measures that gap**

To learn real work, AI needs **a place where real outcomes score every move.**

Sources: Alpha Arena (Nof1, May 2026); Xitadel public report.

Solution

AI acts in the environment; **real outcomes become the next training signal**



Three products, all from the same environments

Environments

Trading and other professional settings rebuilt by real rules, licensed to labs

AI data

Agent trajectories, expert data and evals, produced by the environments

Iterative training

We run training for labs and deliver before/after gains

Expert iteration today; **the goal is AI that improves itself in the real world.**

First result · Xitadel

An open model trained in our trading environment: **up to 12% better**

+12%

Qwen3.8-27B on unseen competition trading days

Best run; mean and confidence interval within 3 months

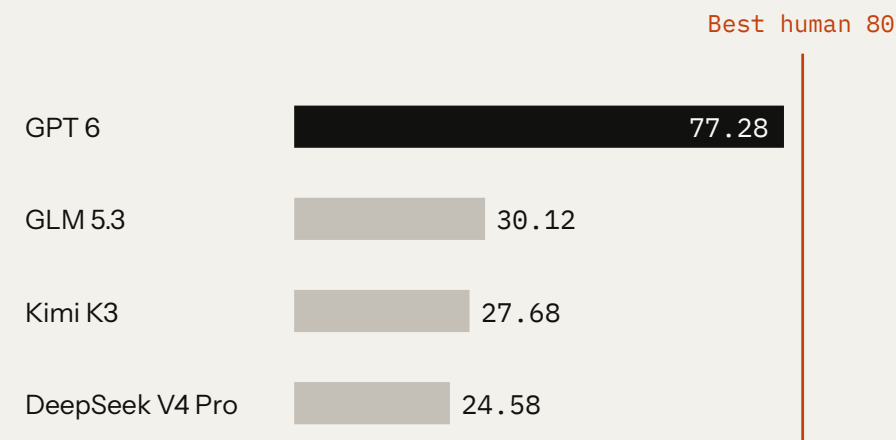
7 trading tasks	Market making, baskets, options, conversion, multi-asset
------------------------	--

Simulator scoring	P&L, drawdown, Sharpe; same strategy, same score
--------------------------	--

Human benchmark 80	Best strategies from 6 competition teams
---------------------------	--

The best model is still below humans: room to train

Xitadel public preview · overall score, 7 tasks



Data: IMC Prosperity competition order books (open source, MIT). Results: Xitadel-QuantBench public report.

Products

7 environments in 14 days, 5 open-sourced, each producing training data

Trading

Xitadel

One training round done: +12%

Open · 102 stars

AI research

SimReal-MLBench

60 research tasks

Open · 178 stars

Prediction

FuturePredict

Settles when real events resolve

Partly open

Math proofs

MathmoBench

Accepts proofs, not guesses

Open · 102 stars

Reasoning

Puzzle Benchmark

749 hand-written puzzles

Open · 17 stars

Finance close

Month-End Close

Passes only with zero errors

Live

Software

SWE-Forward

Code must survive the next version

Live

Every environment yields

Agent trajectories

Expert data · evals

All built with no outside funding

Stars as of September 29, 2026.

Expert network

7,000+ experts signed up from 21 top universities: **the source of our expert data**

7,000+

On the waitlist

Selected universities in the network

21

Top universities

200K+

Reachable



Each domain leaves expert standards and failure cases behind, **reused in the next one**

Network coverage does not mean registration or delivery; logos do not imply endorsement.

Next · Personal agents

Before a personal agent acts for someone, **it practices here**

Muse · Dots

Meta and OpenAI launched personal agents in Sep 2026

→ **Big tech is in; agents now act for people**

How we deliver it

Open

Trading and event prediction

\$10B

Instinct's valuation; its agent books trips and buys

→ **Agents that spend money are a big market**

Pricing

Per practice round, \$2–\$10

Practice first

When an agent errs, the user pays

→ **Practice grounds: our second customer**

Goal

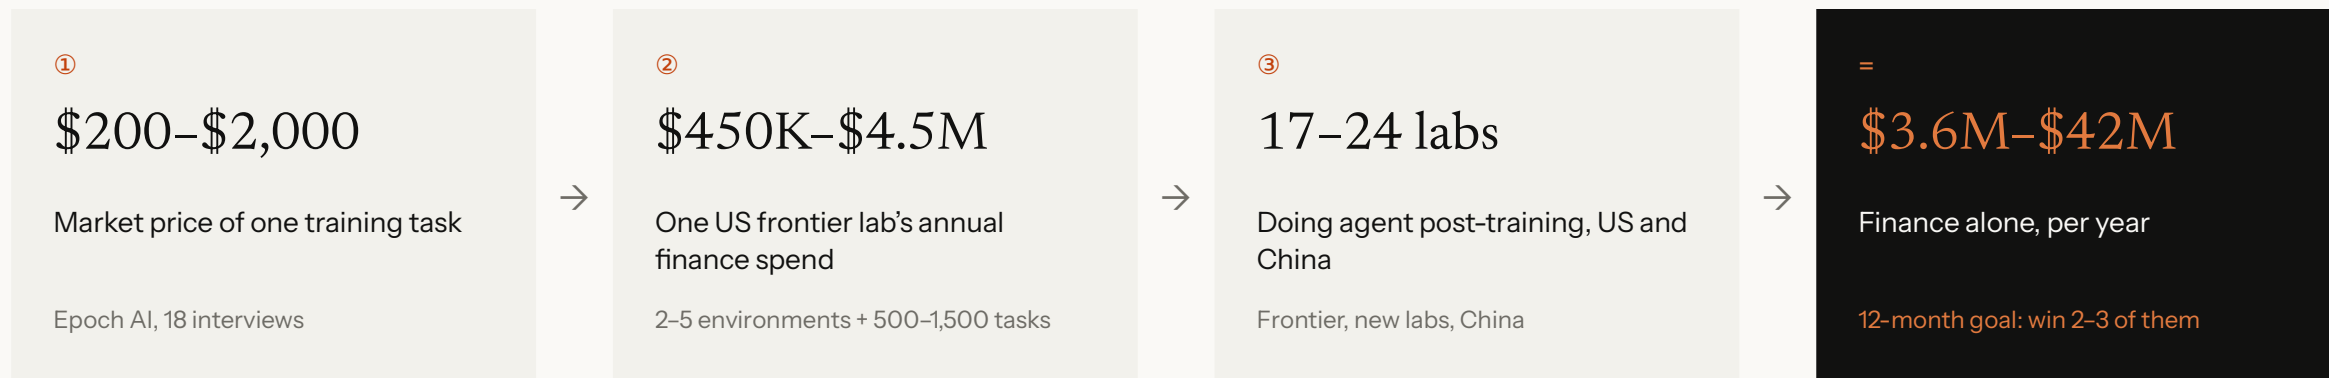
First 10 agent teams within 6 months

De-identified, consented practice runs become lab training data: **same environments, a second customer.**

Sources: Meta and OpenAI launches (Sep 2026); Instinct valuation (Reuters, Sep 28, 2026).

Market

Labs already buy at set prices: **finance alone is about \$3.6M–\$42M a year**



Each new domain adds another

Event prediction and AI research are live; at finance's size, three domains reach ~\$10.8M–\$126M a year

Prices are already set

RL-environment contracts run six to seven figures a quarter; OpenAI pays ex-bankers \$150 an hour for training data

Sources: Epoch AI, An FAQ on RL environments (Jan 2026); Bloomberg (Oct 2025).

Estimates: volume per lab, tier counts (China benchmarked to UniPat orders) and the domain extrapolation are our assumptions, to be set by the first paid pilots.

Competition

Others do one piece; **we run the whole line, from environment to scoring to training**

Who	What they do	What's missing
Expert data platforms Mercor, Surge, Handshake	Expert hours and human grading	Grading by people; not recomputable
Peers in China UniPat, Humanlaya	Evals and post-training data	Mostly domestic labs
AI trading evals Nof1 (Alpha Arena)	Live trading leaderboard	Evaluates, doesn't train
Labs in-house	Domains they already know	Lack rules and experts in specialist fields
SimReal	Environment + scoring + data + training	Trading loop done: +12%

The category is priced: **Bloomberg reports Alibaba plans to lead \$300M in UniPat at ~\$2.5B.**

Why us

Trading's costliest lesson is that backtests lie: **we built it into the scoring**

Tick-level matching

How

Orders queue; each fill is tied to the book at order time

Result

Scores spread out: GPT 6 at 77, three other frontier models at 24–30

Risk in the score

How

P&L minus drawdown, times a Sharpe factor

Result

Losses or drawdowns that eat the profit score 0: two frontier models hit 0 on 2 tasks each

No gaming

How

The last trading day is hidden from the agent; result hashes are published first

Result

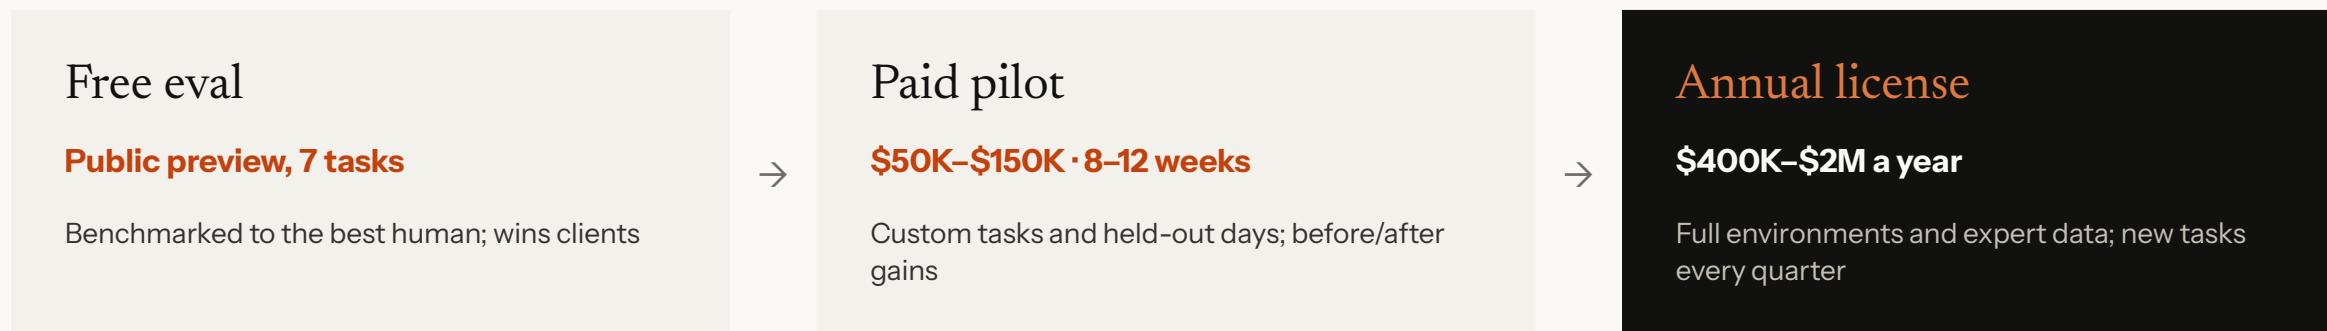
Same strategy, same score; anyone can recompute

Three founders from quant firms; **the team designed these rules itself.**

Source: Xitadel-QuantBench public report and scoring spec (REPORT, SCORING).

Business model & traction

Free evals bring clients in; **paid pilots turn into annual licenses**



2 labs

Frontier labs in purchase talks

Same model: AfterQuery reached a \$100M run-rate in 14 months

Supported by practitioners at trading firms



Logos show where supporters work; they do not imply endorsement.

The round & use of funds

\$6M at \$75M post-money: 2–3 paying labs within 12 months

Raising

\$6M

Seed

Dilution

8%

Pre-money

\$69M

Post-money

\$75M

Revenue engine • \$3M

Environments & data

Environments \$1.05M · data \$600K · delivery \$900K · sales \$450K

3 months

First paid pilot

6 months

Second license

12 months

\$1M–\$3M ARR

R&D engine • \$3M

Iterative training

Compute \$1.35M · research \$1.05M · live trading & compliance \$600K

3 months

Trading gains:
significance test

6 months


Gains in event prediction

12 months

Small live-capital test

Valuation references: Applied Compute \$100M seed post-money (Jun 2025); Mirendil \$1B seed post-money (Jun 2026).

Round: RMB 40M at RMB 500M post-money, at 6.7351 RMB/USD (Sep 30, 2026). Cayman–Hong Kong–onshore structure set up before closing.




Every industry's best AI will come from our worlds.

Three worlds live: trading, AI research, event prediction

SimReal · AI that improves itself in the real world

`business@simreal.co` · `simreal.co` · `x.com/simrealhq`



simreal

Confidential · For invited investors only · September 2026