



AI that improves itself in the real world

Rebuilding the real world from real data


Training ground for AI labs, enterprises, and personal agents

Seed deck

October 2026

Contact

business@simreal.co



Overview

We give AI and personal agents a real world **to fail, learn, and improve themselves**

What we do

Environments & data

For AI labs and enterprises

Proprietary environments for self-improvement

Next: open practice to personal agents

Shipped

7 products, 14 days

Live: trading, AI research, event prediction

5 public repos; no outside funding

Key result

+12%

Qwen3.8-27B after training in Xitadel

Up to 12% better on unseen trading days

Reproduced in multiple independent runs

Team

Quant founders

Born 2005; math at Cambridge, LSE, and Duke

Experience at Jane Street, Citadel,

Optiver, and Millennium

Market

\$8.5B → \$700B

Today: training data and RL environments

2030: company scenario, 82x today

Not an independent forecast

6-month goals

20 customers

3-4 customers in months 1-3; 20 by month 6

RSI on open-source large models by month 6

Today: 2 frontier labs in talks

The best training world makes the best AI. **We start with trading, then the whole world.**

Problem

AI has learned to reason, but it still can't do real-world work

Real work has no answer key, only the world's reaction

	Surface check: passed		Real world: failed
Trading	The backtest only goes up	→	Loses money in markets it hasn't seen
Software	All tests pass	→	Breaks in production
Finance	The books look done	→	Month-end close won't reconcile
Event prediction	The analysis sounds convincing	→	Misses once the outcome is known

6/32

Real-money runs by frontier models that made a profit (Alpha Arena)

77 < 80

On Xitadel, top model GPT 6 scores 77; best human, 80

To learn real work, AI needs a world where real outcomes settle every move.

Why now

The bottleneck is moving from data to environments, **and professional domains still lack good ones**

01

Training has changed

Reasoning models train with RL on automatically checked outcomes

Public human text is projected to run out between 2026 and 2032

→ **Scarce now: auto-scoring environments**

02

Budgets are in place

\$200–\$2,000 per RL task; environment contracts run 6–7 figures a quarter

Mercor acquired RL–environment startup Deeptune in July 2026

→ **Labs now buy environments**

03

Expert tasks: unsolved

Frontier models trading real money made a profit in only 6 of 32 runs

Meta Muse and OpenAI Dots began doing work for people in September 2026

→ **Trading first: easy to verify, big gap**


The window

About 20 RL–environment startups; Wing VC expects 3–5 to last. **Be first to prove gains and win long-term lab contracts**

Sources: Epoch AI (public-text projection; An FAQ on RL environments, Jan 2026); SiliconANGLE, TechCrunch (Jul 2026); Nof1 Alpha Arena (May 2026); Meta and OpenAI launches (Sep 2026); Wing VC (2026). Prices are industry figures, not SimReal contracts.

Solution

Train models on task outcomes, **then check the gains with independent evals**



Self-generated tasks are an R&D direction; eval tasks never feed training or generation

SimReal delivers

Professional data · Training environments · Scoring tools · Eval reports

References: InstructGPT (arXiv 2203.02155); DeepSeek-R1 (arXiv 2501.12948); Absolute Zero (arXiv 2505.03335).

Product

5 benchmarks are public, **and our infrastructure cuts the cost of every training round**


Simreal-MLBench Public

[Public preview] Externally scored agentic ML research benchmark: 60 tasks, real competition ground truth. Open protocol, operated evaluation.

● Python ☆ 252 🏆 32


Xitadel-QuantBench Public

[Public preview] The first public quant trading benchmark scored against what professional traders actually made on the same asset

● Python ☆ 101 🏆 8


MathmoBench Public

[Public preview] An advanced AI Benchmark for math, produced by mathmo at St John's College, Cambridge.

● Python ☆ 101 🏆 8


hard-puzzle-benchmark Public

[Public preview] Hard human-written reasoning puzzles: 749 answered items in 15 topics (707 scored core), with protocol and scorer.

● Python ☆ 44 🏆 4


future-prediction-bench Public

[Public alpha] Probability forecasting protocols, delayed-reward agent rollouts, and verifiable real-world task environments.

● Python ☆ 14 🏆 1

SimReal-MLBench

60 real competition tasks that test ML research skill

Xitadel-QuantBench

Replays order books vs. pro traders on the same asset; keeps trajectories

MathmoBench

Advanced math that tests reasoning and problem solving

Puzzle Benchmark

749 hand-written reasoning puzzles across 15 topics

FuturePredict Bench

Forecasts locked before the deadline, scored on real outcomes

Infrastructure: more training from the same compute

1/4

Time per training round

Deep microVM isolation

+64%

Scored attempts, same resources

Rollout allocation

-1/3

Checkpoint time, half the saves

Smart checkpointing picks when to save

Repos as of Oct 2, 2026. Infra metrics are internal tests with different baselines; they can't be multiplied or added.

Traction

7 products in 14 days with no outside funding; 2 frontier labs in talks

2

Frontier labs in talks

→ Path to our first paid pilot

14

Days to build 7 products

→ How fast we deliver

512

GitHub stars on 5 public repos

→ Dev and researcher interest

\$0

Outside funding

→ Every dollar went into product

As of September 29, 2026; GitHub stars as of October 2, 2026

Practitioners supporting our R&D work at



200,000+

Reachable, verifiable experts

7,000+

Experts on the waitlist

21

Universities: students and alumni



Logos show where supporters work and some of the universities in our network; they do not imply endorsement. Network reach is not registration or participation.

Market

Spend on training data and environments is accelerating: **our 2030 scenario is about 82x today**

Today: training data and RL environments

\$8.5B / year

Estimated revenue of 50+ vendors (Deedy Das, Jul 2026)

2030: company scenario

→ **\$700B / year**

About 82x today; assumes a \$7T AI economy with 10% spent on training

27x

Mercor's gross run-rate: \$75M to \$2B in 16 months

→ **Expert-data demand is growing fast**

10x

Mercor's valuation: \$2B to \$20B (in talks) in 17 months

→ **Investors have priced this category**

18x

Snorkel AI data-services growth in about a year, to a \$375M run-rate

→ **Shift to experts and environments**

Sources: Mercor (TechCrunch, Sacra, Dealroom); Snorkel AI announcement (Sep 2026); Deedy Das industry map (Jul 2026). 2030 is a company scenario, not an independent research forecast.

Competition

We start with trading and event prediction: **others win on scale; we win on scoring and speed**

What clients use today	What they do well	Where they fall short	SimReal's edge
Data & environment vendors Scale AI, AfterQuery, Mercor	Big expert networks, deep lab ties	Mostly general domains; trade matching and risk scoring need desk experience	Desk-designed matching and risk scoring; same strategy, same score
Public benchmarks Alpha Arena and other leaderboards	Easy to compare models	Eval only, no training data; saturate soon after release	Same environment yields training data; reserve tasks keep it fresh
Built in-house	Fits their own tools and business	Lacks domain rules and experts; ties up research staff	Ready-made trading and event-prediction worlds plus experts, on demand

Reusable assets


Tasks, scorers, held-out test sets, expert data, and failure cases grow with each client and model

The category is priced: Bloomberg reports Alibaba plans to lead **UniPat's \$300M round at about a \$2.5B valuation**



Business model

We charge AI labs and enterprises **by delivery scope and usage rights**

Who pays



What they buy



Every model upgrade needs new tasks, trajectories, and tests, **so updates drive repeat purchases**

Pricing reflects the proposed business model, not signed contracts.

Why us

Scoring, training, and delivery must all work at once: **we have results on all three**

Scoring done right

How

Trading-desk rules, tick-level matching
P&L minus drawdown, times a Sharpe factor
Held-out last day; hash published up front

Result

GPT 6 at 77; best human at 80
Other frontier models at 24–30

Training that pays off

How

One environment for evals and training
MicroVM isolation, rollout allocation,
and smart checkpoints cut cost per round

Result

Qwen3.8-27B: up to 12% better
+64% scored attempts, same resources

Fast delivery

How

Our own environment-building tools
200,000+ reachable experts
7,000+ on the waitlist

Result

7 products in 14 days
512 stars on 5 public repos

Harder to copy

Each client leaves tasks, scorers, failure cases, and vetted experts, **reused for the next client and domain**

Sources: Xitadel-QuantBench report and scoring spec (+12%: held-out days of IMC Prosperity order books); infra metrics: internal tests; stars: Oct 2, 2026.

Team

Math, ML, quant trading, and data science; **we turned down top quant firms' return offers to build this**



Charles CEO

First employee at United Stables (U stablecoin)
Helped U go 0 → \$1.4B in a year, Binance in a month
Ran institutional relations; worked with SIG, DRW
Sole intern on HSBC HKD stablecoin; HKMA compliance
X (Twitter) writer with millions of views
Hedge fund intern at Citadel, London

Born 2005 · LSE, Math (undergrad) · SCIE, Shenzhen
USAMO qualifier

Henry CTO

Summer ML research at Cambridge
Supervised by statistician Po-Ling Loh (IMS Fellow)
Youngest undergrad AI researcher, Cambridge center
Quant experience: Jane Street, Citadel, and Optiver
Designed the 5 benchmarks and RL environments
Top 30 worldwide in a Cambridge math contest

Born 2005 · Cambridge Math, First-Class (Scholar)
SCIE, Shenzhen · UK Physics Olympiad Super Gold

Amaris COO

Data scientist at Millennium, Hong Kong
Millennium alt-data team's first-ever new-grad hire
Helped launch Plug and Play's first Hong Kong event
Held with Hong Kong Science Park; 200+ attendees
Host of J&J MedTech's tech summit (200+ attendees)
Published author at 17; 100K+ engagements online

Born 2005 · Duke, Math & Statistics (undergrad)
YK Pao School, Shanghai

Trading-desk discipline, ML research, and institutional ties: **exactly what training environments need**

Next 6 months

The revenue engine sells to today's buyers; **the RSI engine works on self-improvement**

Revenue engine

Environments & data

Licenses and data delivery that build recurring revenue

Focus

Environments and expert data; onboarding, sign-off, and updates

Months 1-3

3-4 customers

Deliver first orders, collect payment

Month 6

20 customers total

Expansion orders and repeat purchases

RSI engine

Recursive self-improvement

R&D on self-generated tasks, training, and evals

Focus

Compute and research; validate in trading and event prediction

Months 1-3

Small-model tests

Validate the self-generated task loop

Month 6

10+ environments

End to end; RSI on open-source LLMs

Self-improving AI in the real world,
from trading to the whole world.

business@simreal.co · simreal.com.cn