




AI that improves itself in the real world

Training environments built from real-world data

For AI labs, enterprises, and personal agents



Company overview

October 2026

Contact

business@simreal.co

At a glance

We build real-world training environments where AI agents **fail, learn, and improve themselves, again and again**

What we do

Environments & data

Agent trajectories, expert data, and evals
For labs and enterprises; personal agents next
Proprietary environments for self-improvement

Team

3 quant founders

Born 2005; math at Cambridge, LSE, and Duke
Experience at Jane Street, Citadel, Optiver,
and Millennium

Shipped

7 products, 14 days

3 worlds live: trading, AI research,
and event prediction
5 public repos; no outside funding

Market

\$8.5B → \$700B

Today: est. revenue, training data + RL envs
2030: company scenario, about 82x today
Not an independent forecast

Key result

+12%

Qwen3.8-27B, after training in Xitadel
Traded up to 12% better on unseen trading days
Reproduced in multiple independent runs

Next 6 months

Goal: 20 customers

3–4 customers in the first 3 months
Month 6: self-improvement on open-source LLMs
Today: 2 frontier labs in talks

Whoever has an industry's best training world has its best AI. **We start with trading, then the whole world.**

Problem

AI has learned to reason, but it still can't do real-world work well

Real work has no answer key, only the world's reaction

	Passes the surface check		Fails in the real world
Trading	The backtest only goes up	→	Loses money in markets it hasn't seen
Software	All tests pass	→	Breaks in production
Finance	The books look finished	→	Month-end accounts don't reconcile
Event prediction	The analysis sounds convincing	→	Misses once the outcome is in

6/32

Alpha Arena (Nof1): 6 of 32 real-money runs profitable

77 < 80

Xitadel: top model GPT 6 scores 77; best human, 80

To learn real work, AI needs a world that scores every move by its real outcome.

Why now

Training's bottleneck is shifting from data to environments: **budgets exist, but professional domains lack good ones**

01

Training has changed

Reasoning models now train with RL on automatically verifiable outcomes

Public human text is projected to run out between 2026 and 2032

→ **Scarce now: auto-scored environments**

02

Budgets are in place

\$200–\$2,000 per RL task; environment contracts run 6–7 figures a quarter

Mercor acquired RL-environment startup Deeptune in July 2026

→ **Labs now buy environments**

03

Pro tasks still unsolved

In Alpha Arena, frontier models trading real money profited in 6 of 32 runs

Meta Muse and OpenAI Dots began doing work for people in September 2026

→ **Trading first: easy to verify, big gap**


The window

About 20 RL-environment startups; Wing VC expects 3–5 to last. **Be first with verifiable gains and long-term lab contracts**

Sources: Epoch AI (public-text projection; RL-environment FAQ, Jan 2026); SiliconANGLE, TechCrunch (Jul 2026); Nof1 (May 2026); Meta, OpenAI (Sep 2026); Wing VC (2026). Muse, Dots, and task prices are industry examples, not SimReal customers or contracts.

Solution

We train models on task outcomes **and check the gains with independent evals**



Self-generated tasks are an R&D direction; eval tasks never feed training or new tasks

SimReal delivers

Professional data · Training environments · Scoring tools · Eval reports

References: InstructGPT (arXiv 2203.02155); DeepSeek-R1 (arXiv 2501.12948); Absolute Zero (arXiv 2505.03335).

Product

5 benchmarks are public; **our training infrastructure cuts the cost of every round**


Simreal-MLBench Public

[Public preview] Externally scored agentic ML research benchmark: 60 tasks, real competition ground truth. Open protocol, operated evaluation.

● Python ☆ 252 🏆 32


Xitadel-QuantBench Public

[Public preview] The first public quant trading benchmark scored against what professional traders actually made on the same asset

● Python ☆ 101 🏆 8


MathmoBench Public

[Public preview] An advanced AI Benchmark for math, produced by mathmo at St John's College, Cambridge.

● Python ☆ 101 🏆 8


hard-puzzle-benchmark Public

[Public preview] Hard human-written reasoning puzzles: 749 answered items in 15 topics (707 scored core), with protocol and scorer.

● Python ☆ 44 🏆 4


future-prediction-bench Public

[Public alpha] Probability forecasting protocols, delayed-reward agent rollouts, and verifiable real-world task environments.

● Python ☆ 14 🏆 1

SimReal-MLBench

60 real competition tasks that test ML research ability

Xitadel-QuantBench

Replays order books, scores vs. pro traders, keeps trajectories

MathmoBench

Advanced math problems that test reasoning and problem solving

Puzzle Benchmark

749 hand-written reasoning puzzles across 15 topics

FuturePredict Bench

Locks forecasts and evidence by the deadline; scored on outcomes

Infrastructure: more training from the same compute

1/4

Time per training round, vs. baseline
Deep microVM isolation

+64%

Scored attempts, same resources
Rollout allocation

-1/3

Checkpoint time
Smart checkpointing; half as many saves

Public repos as of Oct 2, 2026. Infra metrics are internal tests on different baselines; they can't be multiplied or added.

Traction

7 products in 14 days with no outside funding; 2 frontier labs in talks and 200,000+ experts reachable

2

Frontier labs in talks

→ Path to our first paid pilot

14

Days to build 7 products

→ How fast we deliver

512

GitHub stars, 5 public repos

→ Dev and researcher interest

\$0

Outside funding

→ Every dollar went into product

As of September 29, 2026; GitHub stars as of October 2, 2026

Practitioners supporting our R&D work at



200,000+

Reachable, verifiable experts

7,000+

Experts on the waitlist

21

Universities: students and alumni



Logos show where supporters work and some of the universities in our network; they do not imply endorsement. Network reach is not registration or delivery.

Market

Spending on training data and environments is accelerating: **our own 2030 scenario is about 82x today**

Today: training data and RL environments

\$8.5B / year

Estimated revenue of 50+ vendors (Deedy Das, Jul 2026)

2030: company scenario estimate

→ **\$700B / year**

About 82x today; assumes a \$7T AI economy with 10% spent on training

27x

Mercor's gross run-rate grew from \$75M to \$2B in 16 months

→ **Demand for expert data is surging**

10x

Mercor's valuation: \$2B to \$20B in 17 months (\$20B in talks)

→ **Capital has priced this category**

18x

Snorkel AI data services, in about a year; \$375M run-rate

→ **Shift to experts and environments**

Sources: Mercor (TechCrunch, Sacra, Dealroom); Snorkel AI announcement (Sep 2026); Deedy Das industry map (Jul 2026). 2030 is a company scenario estimate, not an independent research forecast.

Competition

We start with trading and event prediction: **others win on scale; we win on domain scoring and speed**

What clients use today	What they do well	Where they fall short	SimReal's edge
Data & environment vendors Scale AI, AfterQuery, and Mercor	Large expert networks, deep lab relationships	Mostly general domains; order matching and risk scoring need desk experience	Desk-built order matching and risk scoring; same strategy, same score
Public benchmarks & evals Alpha Arena and other leaderboards	Easy to compare models	Evals only, no training data; saturate soon after release	Same environment yields training data; reserve tasks keep tests fresh
Built in-house The client's own team	Fits their own tools and business	Lacks domain rules and experts; ties up research headcount	Ready-made trading and event-prediction environments, plus experts

Reusable assets


Tasks, scorers, held-out test sets, expert data, and failure cases grow with each client and model

Already priced: Bloomberg reports Alibaba plans to lead a **\$300M round in UniPat** at about a **\$2.5B valuation**



Business model

We charge AI labs and enterprises **by delivery scope and usage rights**

Who pays



What they buy



Every model upgrade needs new tasks, trajectories, and tests, **so ongoing updates bring repeat purchases**

Pricing reflects the proposed business model, not signed contracts.

Why us

Scoring, training, and delivery must all work at once: **we have results on all three**

Scoring done right

How

Trading-desk rules: tick-level matching
P&L minus drawdown, times Sharpe
Last day held out; hashes pre-published

Result

GPT 6 at 77; best human at 80
Other frontier models at 24–30

Training that pays off

How

One environment for evals and training
MicroVM isolation, rollout allocation,
and smart checkpointing cut cost per round

Result

Qwen3.8-27B: up to +12% on unseen days
+64% scored attempts, same resources

Fast delivery

How

Our own environment-building tools
200,000+ reachable experts
7,000+ experts on the waitlist

Result

7 products, 14 days, no outside funding
512 stars across 5 public repos

Harder to copy

Each client will leave tasks, scorers, failure cases, vetted experts, **reused for the next client and the next domain**

Sources: Xitadel-QuantBench report and scoring spec (IMC Prosperity order books; +12% on unseen days); infra: internal tests; stars as of Oct 2, 2026.

Team

Math, ML, quant trading, and data science: **we turned down top quant firms' return offers to build this**



Charles CEO

First employee at United Stables (U stablecoin)
Helped U: 0 → \$1.4B in a year, on Binance in a month
Ran institutional relations; worked with SIG, DRW
Sole intern on HSBC HKD stablecoin; HKMA compliance
Hedge fund intern at Citadel, London
X (Twitter) writer with millions of views

Born 2005 · LSE, Math (undergraduate) · USAMO qualifier
Shenzhen College of International Education (SCIE)

Henry CTO

Summer ML research at Cambridge
Mentored by statistician Po-Ling Loh, IMS Fellow
Youngest undergrad AI researcher, Cambridge center
Quant experience: Jane Street, Citadel, and Optiver
Designed our 5 benchmarks and RL environments
Top 30 worldwide in a Cambridge math contest

Born 2005 · Cambridge, Math, First-Class (Scholar)
SCIE, Shenzhen · UK Physics Olympiad Super Gold

Amaris COO

Data scientist at Millennium Hong Kong
First-ever new-grad hire, Millennium alt-data team
Helped launch Plug and Play's first HK event
(with Hong Kong Science Park; 200+ attendees)
Host, J&J MedTech tech summit (200+ attendees)
Published author at 17; 100K+ engagements online

Born 2005 · Duke, Math & Statistics (undergraduate)
YK Pao School, Shanghai

Trading-desk discipline, ML research, and institutional ties: **the three things professional environments need**

Next 6 months

The revenue engine sells to today's buyers; **the RSI engine works on recursive self-improvement**

Revenue engine

Environments & data

Licenses and data delivery for recurring revenue

Focus

Environments and expert data; onboarding, acceptance, updates

Months 1-3

3-4 customers

Deliver first orders and collect payment

Month 6

20 customers in total

Expansion orders and repeat purchases

RSI engine

Self-improvement

R&D: self-generated tasks, training, independent evals

Focus

Compute and research; validate in trading and event prediction

Months 1-3

Small-model tests

Validate the self-generated task loop

Month 6

10+ environments

End to end; RSI on open-source LLMs

Self-improving AI in the real world,
from trading to the whole world.

business@simreal.co · simreal.com.cn