




AI that improves itself in the real world

RL training environments and AI data

Quant founders born 2005. 7 products in 14 days.



Company overview

October 2026

Contact

business@simreal.co

Overview

RL training environments for AI labs in professional domains, **starting with trading**

Team

Quant founders

Math at Cambridge, LSE, and Duke
Jane Street, Citadel, and other quant firms

Speed

7 products, 14 days

No outside funding
5 public repos, 512 stars

Result

+12%

Qwen3.8-27B trained in Xitadel
Up to +12% on unseen trading days; reproduced multiple times

What we do

Environments & data

RL environments, expert data, agent data
Live: trading, AI research, event prediction

Market

\$8.5B

Est. annual revenue of data and RL-env vendors
Mercor gross run-rate up 27x in 16 months

Next

Personal agents

Practice environment in development
Today: 2 frontier labs in talks

Team

Three founders born 2005, all from quant firms; we turned down return offers to build SimReal



Charles CEO

First employee at United Stables (U stablecoin)
 Helped U: 0 → \$1.4B in a year, on Binance in a month
 Ran institutional relations; worked with SIG, DRW
 Sole intern on HSBC HKD stablecoin; HKMA compliance
 Hedge fund intern at Citadel, London
 X (Twitter) writer with millions of views

Born 2005 · LSE, Math (undergraduate) · USAMO qualifier
 Shenzhen College of International Education (SCIE)

Henry CTO

Summer ML research at Cambridge
 Mentored by statistician Po-Ling Loh, IMS Fellow
 Youngest undergrad AI researcher, Cambridge center
 Quant experience: Jane Street, Citadel, and Optiver
 Designed our 5 benchmarks and RL environments
 Top 30 worldwide in a Cambridge math contest

Born 2005 · Cambridge, Math, First-Class (Scholar)
 SCIE, Shenzhen · UK Physics Olympiad Super Gold

Amaris COO

Data scientist at Millennium Hong Kong
 First-ever new-grad hire, Millennium alt-data team
 Helped organize Plug and Play's first HK event
 (with Hong Kong Science Park; 200+ attendees)
 Host, J&J MedTech tech summit (200+ attendees)
 Published author at 17; 100K+ engagements online

Born 2005 · Duke, Math & Statistics (undergraduate)
 YK Pao School, Shanghai

Speed

No outside funding. 7 products in 14 days, across 7 domains

Speed

14 days

7 products, 7 domains

→ 5 open source, 2 private

Acceleration

1/4

Time per training round

→ Internal test: more training per GPU

Expert bench

7,000+

Experts on the waitlist

→ Signed up; vetted by domain

Our own build tools

Built in-house; task libraries and scorers grow with each delivery

Training infrastructure

+64% scored attempts on the same resources (internal test)

Acceptance before delivery

Tested on unseen tasks; delivered only at the agreed bar

After the round

Expert data

Vet the waitlist; trading experts first

RL environments

License-ready trading and forecasting environments

Agent data

Personal-agent practice environment in beta

7 products in 14 days as of Sep 29, 2026 (5 open source, 2 private) across trading, AI research, event prediction, math, reasoning, month-end close, and software engineering. Infra metrics are internal tests with different baselines; they can't be multiplied or added. Post-round steps are plans.

Traction

5 open benchmarks, 512 stars; 2 frontier labs in talks

2

Frontier labs in talks

→ Toward a first paid pilot

7

Products in 14 days

→ 3 live: trading, AI research, forecasting

512

GitHub stars

→ 5 public repos

\$0

Outside funding

As of September 29, 2026; GitHub stars as of October 2, 2026

Practitioners supporting our R&D work at



200,000+

Reachable, verifiable experts

7,000+

Experts on the waitlist

21

Universities: students, alumni



Logos show where supporters work and some universities in our network; they do not imply endorsement. Network reach is not registration or delivery.

Problem

Models can reason, **but fail at real-world work**

	Looks like		Actually
Trading	Backtest only goes up	→	Loses money in unseen markets
Software	All tests pass	→	Breaks in production
Finance	Books look done	→	Month-end won't reconcile
Forecasting	Analysis sounds right	→	Calls the outcome wrong

6/32

Frontier models trading real money: 6 of 32 runs made a profit (Alpha Arena)

Labs lack **training environments scored by real outcomes.**

Why now

Labs now pay for training environments; **supply in professional domains is short**

01

Training changed

Reasoning models train with RL on automatically verifiable outcomes

Public human text projected to run out 2026–2032

→ **Auto-scored environments are scarce**

02

Money is flowing

\$200–\$2,000 per RL task

Environment contracts: 6–7 figures a quarter

Mercor acquired environment startup Deeptune

→ **Labs now buy environments**

03

Agents launched

Meta Muse and OpenAI Dots

Launched Sep 2026; doing work for people

→ **The next set of buyers**


Window

About 20 early RL-environment startups; Wing VC expects 3–5 to last. **We start with trading, where results are checked automatically.**

Sources: Epoch AI (public-text projection; An FAQ on RL environments, Jan 2026); SiliconANGLE, TechCrunch (Jul 2026); Meta, OpenAI launches (Sep 2026); Wing VC (2026). Muse and Dots are industry examples, not SimReal customers.

Solution

Train on task outcomes, **accept on tasks the model has never seen**



Acceptance tasks never feed training or task generation

We deliver

Training environments · Expert data · Scoring tools · Eval reports

References: InstructGPT (arXiv 2203.02155); DeepSeek-R1 (arXiv 2501.12948); Absolute Zero (arXiv 2505.03335).

Product

5 benchmarks are open source; the eval environments double as training environments


Simreal-MLBench Public

[Public preview] Externally scored agentic ML research benchmark: 60 tasks, real competition ground truth. Open protocol, operated evaluation.

● Python ☆ 252 🏆 32


Xitadel-QuantBench Public

[Public preview] The first public quant trading benchmark scored against what professional traders actually made on the same asset

● Python ☆ 101 🏆 8


MathmoBench Public

[Public preview] An advanced AI Benchmark for math, produced by mathmo at St John's College, Cambridge.

● Python ☆ 101 🏆 8


hard-puzzle-benchmark Public

[Public preview] Hard human-written reasoning puzzles: 749 answered items in 15 topics (707 scored core), with protocol and scorer.

● Python ☆ 44 🏆 4


future-prediction-bench Public

[Public alpha] Probability forecasting protocols, delayed-reward agent rollouts, and verifiable real-world task environments.

● Python ☆ 14 🏆 1

SimReal-MLBench

60 real competition tasks that test ML research

Xitadel-QuantBench

Replays order books against pro traders on the same asset; keeps trajectories

MathmoBench

Advanced math that tests reasoning and problem solving

Puzzle Benchmark

749 hand-written reasoning puzzles across 15 topics

FuturePredict Bench

Locks forecasts and evidence before the deadline; scored on real outcomes

Training infrastructure

1/4

Time per round
microVM isolation

+64%

Scored attempts, same resources
Rollout allocation

-1/3

Checkpoint time, half the saves
Smart checkpointing

Public repos as of Oct 2, 2026. Infra metrics are internal tests with different baselines; they can't be multiplied or added.

Personal agents

Personal agents have launched; **our practice environment is in development**

Meta Muse and OpenAI Dots launched in Sep 2026 to do work for people. Agents need practice before they act.

01

Real app replicas

Email, calendar, and docs
in VMs and browsers

02

Long cross-app tasks

Multi-step tasks
scored on the final result

03

Failure recovery

Pop-ups, dropped connections,
expired logins: can it recover?

Who pays

Personal-agent teams pay for practice; **labs buy the trajectories it produces**

In development; features are planned. Muse and Dots are industry examples, not SimReal customers.

Market

\$8.5B today; **top vendors grew revenue 10x+ in 16 months**

Today: training data and RL environments

\$8.5B / yr

Est. revenue of 50+ vendors (Deedy Das, Jul 2026)

2030: company scenario

→ **\$700B / yr**

Assumes a \$7T AI economy, 10% spent on training

27x

Mercor gross run-rate: \$75M to \$2B in 16 months

→ **Expert-data demand is rising**

10x

Mercor valuation: \$2B to \$20B in 17 months (in talks)

→ **Capital is doubling down**

18x

Snorkel AI new data services in about a year; \$375M run-rate

→ **Demand shifts to experts and environments**

Comparable Bloomberg: Alibaba plans to lead a \$300M round in UniPat at about \$2.5B (Sep 2026)

Sources: Mercor (TechCrunch, Sacra, Dealroom); Snorkel AI announcement (Sep 2026); Deedy Das industry map (Jul 2026); UniPat (Bloomberg, Sep 2026). 2030 is a company scenario, not an independent forecast.

Business model

Sell to labs first, **then enterprises and personal-agent teams**

Who pays

Now

AI labs

2 frontier labs in talks

Expanding

Enterprise agent teams

Train, test, accept before launch

Next

Personal-agent teams

Long, cross-app tasks; recovery

What they buy

Environments

Executable tasks and scoring tools

Pricing By term, scope, usage rights

AI data

Agent trajectories, expert and eval data

Pricing By volume or batch

Joint training

Experiments, transfer tests, managed runs

Pricing Project or ongoing contract

Scope & acceptance



Paid pilot



Licensed delivery



Ongoing updates

Every model upgrade needs new tasks and test sets, **so labs keep buying**

Pricing reflects the proposed business model, not signed contracts.

Competition

Others win on scale; we win on domain scoring and speed. **Scale here means labs won**

Alternatives	Strengths	Gaps	SimReal
Data & env vendors Scale AI, AfterQuery, Mercor	Large, many experts Deep lab ties	Human-heavy; grow by headcount Trade-level scoring isn't core	Quant-firm team; desk-rule scoring 7 products in 14 days
Public benchmarks Alpha Arena and others	Easy to compare models	Eval only, no training data Saturate soon after release	Same environment yields data Reserve tasks keep it fresh
Built in-house	Fits their own work	Lacks domain rules and experts Ties up research staff	Trading, forecasting envs ready 200,000+ experts reachable

The scale gap

Over 90% of Mercor’s H1 2026 gross revenue came from three labs.
We have 2 in talks; the round adds delivery and compute.

Mercor revenue concentration: The Information (Jul 2026): about \$614M gross revenue in H1 2026, over 90% from three foundation-model companies. Competitor names are industry examples.

Strategy & investment

The round funds three lines: **expert data**, **RL environments**, **agent data**

Use of funds Revenue engine 50%: production, delivery, sales | RSI engine 50%: compute, research, validation

Expert data

Benchmark Mercor: \$2B gross run-rate (Jun 2026)

Now

200,000+ reachable experts; 7,000+ on the waitlist

Investment

3 delivery leads: vetting, QA, acceptance
Vet the waitlist; trading and forecasting first
Two-reviewer checks, agreement rate reported

Next

Vetted experts in 7 domains, sold by batch

RL environments

Benchmark AfterQuery: \$100M run-rate in 14 months

Now

7 domains built; trading, AI research, forecasting live

Investment

3 env engineers: licensable trading, forecasting
3 researchers + compute: self-generated tasks
Red-teaming against reward hacking

Next

20+ licensable environments; pilots become licenses

Agent data

Benchmark Meta, OpenAI personal agents (Sep 2026)

Now

Xitadel keeps trajectories; practice environment in development

Investment

2 env engineers: email, calendar, docs replicas
Batch trajectories, expert spot checks
2 BD hires: labs, enterprises, agent teams

Next

Personal-agent teams onboard; trajectories sold to labs

Rivals took 14–16 months to scale revenue; **we lead on scoring and speed**

Use of funds, headcount, and next steps are plans. Benchmarks: Mercor (Forbes, Dealroom, Jun 2026); AfterQuery (Business Wire, Apr 2026); Meta, OpenAI launches (Sep 2026).