



AI that improves itself in the real world


Real data. Real environments. AI that gets it right.

Founded

September 10, 2026

Business plan

October 2026



Overview

SimReal is a neolab building **data, RL environments and recursive self-improvement (RSI)** for next-generation AI

Revenue

\$6.75M ARR

24 days after founding
No outside capital

Open source

602 GitHub stars

7 products shipped in 14 days
5 of them open benchmarks

Research result

+12%

Training loop working in trading by week 2
Up to +12% on unseen trading days

Team

Gen-Z quant team

Cambridge, LSE, Duke mathematics
Jane Street, Citadel, Optiver, Millennium

Market

\$8.5B

Annual spend on training data and RL envs
Mercor revenue up 27x in 16 months

Network

200K+ experts

7,000+ more on the waitlist
Students and alumni of 21 universities

Team

Three founders, all 21. They turned down \$0.9M+ in first-year pay to build the next training engine



Charles CEO 21

First employee at United Stables: took USDU from 0 to \$1.4B in a year; Binance listing in a month

Led institutional partnerships (SIG, DRW)

Sole intern on HSBC HKD stablecoin project; worked on HKMA compliance

Citadel hedge fund intern, London

LSE Mathematics · SCIE
USAMO qualifier

Henry CTO 21

Cambridge Mathematics, First Class, scholar

Cambridge ML research with Po-Ling Loh (IMS Fellow)

Quant at Jane Street, Citadel, Optiver

Designed all 5 open benchmarks

SCIE
Top 30 worldwide, Cambridge maths contest · BPhO
Super Gold

Amaris COO 21

Data scientist, Millennium Hong Kong

First graduate hire, Millennium alt-data team

Co-organized Plug and Play's first HK event (with HKSTP, 200+ attendees)

Duke Mathematics & Statistics · YCIS Shanghai

Problem

AI can reason. **It still fails at real work.**

	Looks like		Actually
Trading	Backtest climbs	→	Loses money live
Software	All tests pass	→	Breaks in production
Finance	Books look closed	→	Month-end won't reconcile
Forecasting	Analysis sounds right	→	Outcome proves it wrong

6/32

Frontier models trading real money made a profit in only 6 of 32 runs (Alpha Arena)

What is missing: **a training ground scored by real outcomes.**

Why now

The bottleneck moved from data to environments. **Budgets exist; expert-domain environments don't**

01

Training changed

Reasoning models learn by RL on auto-graded tasks

Public human text runs out 2026–2032

→ **Auto-graded environments are scarce**

02

Money is moving

\$200–2,000 per RL task

Mercor bought env company Deeptune (Jul 2026)

RSI neolab Recursive valued at \$4.65B

→ **Labs and capital pay for envs and RSI**

03

Expert work unsolved

Models still fail at trading and finance

Meta, OpenAI launched personal agents (Sep)

→ **Trading is clear-cut: the place to start**

Window

~20 early RL-environment companies; expect 3–5 to remain. **First to show gains wins the long contracts**

Sources: Epoch AI (Jan 2026); TechCrunch (Jul 2026); SiliconANGLE (May 2026); Meta and OpenAI announcements (Sep 2026); Wing VC (2026).

Market

\$8.5B today. The leader grew revenue 27x in 16 months

Today: training data and RL environments

\$8.5B / year

Combined revenue of 50+ vendors

2030: company estimate

→ **\$700B / year**

A \$7T AI economy, 10% spent on training

27x

Mercor run-rate revenue,
\$75M to \$2B in 16 months

→ **Expert-data demand is exploding**

10x

Mercor valuation,
\$2B to \$20B in 17 months (in talks)

→ **Capital is doubling down**

18x

Snorkel AI new data business,
growth in one year

→ **Demand shifts to experts and envs**

Sources: Deedy Das market map (Jul 2026); Mercor (TechCrunch, Dealroom); Snorkel AI (Sep 2026). 2030 figure is a company estimate.

Customers

AI labs first, **then enterprises and consumers**

Now

AI labs

What they need

Environments and data
to train stronger models

What we provide

Env licenses, training data, evals

Next

Enterprise agent teams

What they need

Turn workflows into practice tasks;
train and test before go-live

What we provide

Custom envs and acceptance evals

Then • B2C

Consumers

What they need

AI that knows them
and fits their day


What we provide

AI tuned to personal habits, fast

One set of environments, **three customer groups**

RSI method

Attempt, get scored, fix the weak spots: **stronger every loop**



Where we are stronger

Expert scoring

Trading-desk standards: top model 77, rest 24-30

Fast, precise tasks

Reviewed by 200K+ experts

Verified gain

+12%

Up to +12% on
unseen trading days

Scores from the public Xitadel benchmark. +12% is the best of several independent replications.

Product

7 products in 14 days. 5 open benchmarks, 602 GitHub stars

602

GitHub stars

5 open benchmarks, as of Oct 4, 2026

Pinned

[Customize pins](#)

 **Simreal-MLBench** Public

[Public preview] Externally scored agentic ML research benchmark: 60 tasks, real competition ground truth. Open protocol, operated evaluation.

Python 302 40

 **Xitadel-QuantBench** Public

[Public preview] The first public quant trading benchmark scored against what professional traders actually made on the same asset

Python 102 8

 **MathmoBench** Public

[Public preview] An advanced AI Benchmark for math, produced by mathmo at St John's College, Cambridge.

Python 101 8

 **hard-puzzle-benchmark** Public

[Public preview] Hard human-written reasoning puzzles: 749 answered items in 15 topics (707 scored core), with protocol and scorer.

Python 70 6

 **future-prediction-bench** Public

[Public alpha] Probability forecasting protocols, delayed-reward agent rollouts, and verifiable real-world task environments.

Python 27 2

SimReal-MLBench **302★**

60 real ML competitions: can AI do research?

Xitadel-QuantBench **102★**

AI trading, scored against the best human strategy

MathmoBench **101★**

Hard mathematics

Puzzle Benchmark **70★**

749 human-written reasoning puzzles

FuturePredict **27★**

Forecast real events, scored on the outcome

Live training environments

Trading · AI research · Forecasting

Why us

Three shifts in RSI in 2026. **We already deliver on each**

What changed in 2026		SimReal
AI does research OpenAI: 3.1 agent-workdays per researcher-workday OpenAI, Sep 2026	→	MLBench: the exam for AI research 302 GitHub stars
AI improves itself An agent rewrote its own code: 7 gains in 8 days Weco AIIDE², Sep 2026	→	RSI working in trading Up to +12% on unseen days
Environments are the bottleneck Noisy env: top agents fall from 83.9% to 57.6% Breaking the Environment Wall, Sep 2026	→	200K+ experts Millions of tasks delivered in 3 weeks

So RSI is gated by environments, our core research. **Most neolabs have no revenue; we hit \$6.75M ARR in 24 days**

Sources: OpenAI (Sep 6, 2026); Weco AI, arXiv 2609.26457; Breaking the Environment Wall, arXiv 2609.29773; Deedy Das neolab list (May 2026). +12% is the best of several replications.

Expert network

Practitioners in our R&D **come from top trading firms**



200K+

Reachable experts

7,000+

Experts on the waitlist

Write · Score · Review

What experts do in our envs

Logos show where practitioners work; no endorsement implied.

University network

Students and alumni of 21 top universities



21

Universities in our student and alumni network

From entry-level roles to frontier research

Selected universities shown; no endorsement implied.

Competition

They win on scale. **We win on expert scoring and speed**

Today's option	Strength	Weakness	SimReal
Data vendors Scale AI, Mercur	Headcount, lab relationships	Labor-bound; no expert scoring	Quant-grade scoring
Public leaderboards Alpha Arena	Easy model comparison	Eval only; saturates fast	Same environment evals and trains
Labs build in-house	Fits their own needs	No domain rules or experts; costs researchers	Ready environments, 200K+ experts

They scale headcount; we scale speed: **7 products in 14 days, \$6.75M ARR in 24**

Business model

Licenses, data, joint training. **Every model upgrade means a repeat purchase**

01

Environment licenses

Trading, AI research, forecasting environments

Priced by term and scope

02

Data

Every run produces data: trajectories, expert demos, eval sets

Priced by volume

03

Joint training

We train customer models in our environments

Priced per project

Engagement

Scope and acceptance



Paid pilot



License and delivery




Ongoing refresh

Each model upgrade needs new tasks: **environments and data refresh, customers keep paying**

Traction

Founded September 10, 2026. **\$6.75M ARR in 24 days**



Infrastructure figures are internal tests on different bases; do not multiply or add them.