

Every industry's best AI will come from our worlds


An AI-native neolab building RL environments, verifiers and training data for agents that improve themselves

Founded

September 10, 2026

Investor presentation

October 2026



Overview

An AI-native neolab for recursive self-improvement (RSI): **the environments, verifiers and data that agents learn from**

Revenue

\$7M ARR

Reached 24 days after founding
No outside capital

Product

7 environments

Trading, ML research, math, reasoning,
forecasting, software and accounting

Proof

+12%

A model trained in our trading environment
did better on trading days it had never seen

Open source

602 GitHub stars

Across 5 open benchmarks,
all shipped in our first 14 days

Team

Quant founders, all 21

Cambridge, LSE and Duke mathematics
Jane Street, Citadel, Optiver, Millennium

Market

\$8.5B a year

Spent today on training data
and RL environments

Team

Three founders, all 21. Together they turned down \$0.9M+ in first-year pay to build SimReal



Charles CEO · 21

First employee at United Stables, where \$U grew from zero to \$1.4B in a year and listed on Binance within a month

Led institutional partnerships with YZi Labs, SIG, DRW and Wintermute

Only intern on HSBC's HKD stablecoin project, working on HKMA compliance

Hedge fund intern, Citadel London

LSE, Mathematics
USAMO qualifier

Henry CTO · 21

Cambridge Mathematics, First Class; scholar of St John's College

Youngest researcher at the Cambridge AI research center, under Prof. Po-Ling Loh (IMS Fellow)

Ex-quant at Jane Street, Citadel and Optiver

Designed all 5 open benchmarks

Top 30 worldwide, Cambridge maths competition
British Physics Olympiad Super Gold

Amaris COO · 21

Data scientist, Millennium Hong Kong

First graduate hire on Millennium's alternative-data team

Co-organized Plug and Play's first Hong Kong event, with HKSTP (200+ attendees)

Duke, Mathematics and Statistics

Problem

Agents can reason. **They still fail at real work**

	Looks right		What actually happens
Trading	Backtest climbs	→	Loses money live
Software	All tests pass	→	Breaks in production
Accounting	Books look closed	→	Month-end doesn't reconcile
Forecasting	Analysis sounds right	→	Outcome proves it wrong

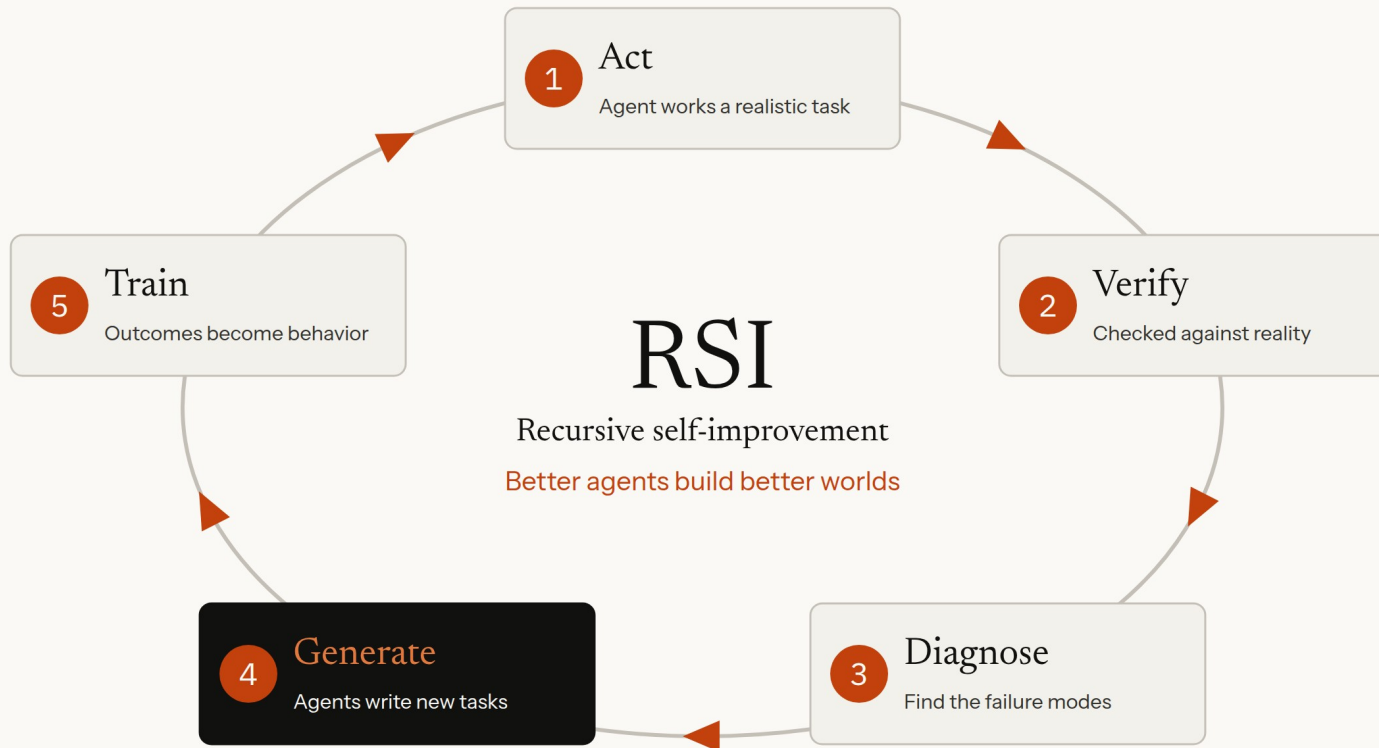
6/32

Frontier models trading real money were profitable in only 6 of 32 runs (Alpha Arena)

The gap: **agents have nowhere realistic to practice, and no source of truth to learn from**

Solution

A loop that turns real outcomes into better agents, **with AI inside it from day one**



What's different

Today

Humans hand-write every task, label and rubric

SimReal

Agents help write new tasks and data, find failure modes and improve the environments they train in. Next: the verifiers and rubrics themselves

Proof point

+12%

on unseen trading days
after training in Xitadel

+12% is the best of several independent replications, measured on trading days held out from training.

Product

Seven environments in 14 days. Each is scored by something the agent can't argue with

602

**GitHub stars across
5 open benchmarks**

As of October 4, 2026

Give the agent a world to practice in.


Make the reward come from a market, a proof, hidden labels, an executable system or a real-world outcome.

Domain	Environment	Reward comes from	Stars
Trading	Xitadel-QuantBench	The market: replayed on unseen days, vs. the best human	102
ML research	SimReal-MLBench	Hidden labels: official competition scoring	302
Math	MathmoBench	A proof	101
Reasoning	Puzzle Benchmark	Hidden answers to 749 human-written puzzles	70
Forecasting	FuturePredict	A real-world outcome, after the deadline	27
Software	SWE-Forward	An executable system: hidden tests	Private
Accounting	Month-End Close	The books must reconcile	Private

SWE-Forward and Month-End Close are not open-sourced.

Traction

Founded September 10, 2026. **\$7M ARR in 24 days, with no outside capital**



Infrastructure figures are internal tests on different bases; do not multiply or add them.

Why now

Agents are starting to improve themselves. **Environments are now the bottleneck**

What changed in 2026		What we have built
AI is doing AI research OpenAI: 3.1 agent-workdays for every researcher-workday OpenAI, Sep 2026	→	SimReal-MLBench: an exam for AI research 302 GitHub stars
AI is improving itself An agent rewrote its own code: 7 gains in 8 days Weco AIDE², Sep 2026	→	RSI working in trading Up to +12% on unseen trading days
Environments are the bottleneck In a noisy environment, top agents fall from 83.9% to 57.6% Breaking the Environment Wall, Sep 2026	→	Seven environments in 14 days 200K+ experts to build more

Our position

Environments gate RSI; they are our core research. **Most neolabs have no revenue; we hit \$7M ARR in 24 days**

Sources: OpenAI (Sep 6, 2026); Weco AI, arXiv 2609.26457; Breaking the Environment Wall, arXiv 2609.29773; Deedy Das neolab list (May 2026). +12% is the best of several replications.

Market

\$8.5B a year today. **One vendor, Mercor, grew revenue 27x in 16 months**

Today: training data and RL environments

\$8.5B / year

Combined revenue of 50+ vendors

2030: the AI economy

→ **\$700B / year**

If AI becomes a \$7T economy and 10% goes to training

27x

Mercor run-rate revenue:
\$75M to \$2B in 16 months

→ **Demand for expert data is exploding**

10x

Mercor valuation:
\$2B to \$20B in 17 months (in talks)

→ **Capital is doubling down**

18x

Snorkel AI new data business:
growth in one year

→ **Demand is shifting to experts and environments**

Sources: Deedy Das market map (Jul 2026); Mercor (TechCrunch, Dealroom); Snorkel AI (Sep 2026). The 2030 figure is our estimate.

Business model

Data is the revenue engine. **Frontier labs today; enterprise and personal agents next**

Now

Frontier AI labs

What they need

Environments, verified trajectories and training data for post-training

What we provide

Environment licenses, data and evals

Next

Enterprise agents

What they need

Vertical AI companies that need agents to be reliable at one job before they deploy

What we provide

Custom environments and acceptance tests

Later

Personal agents

What they need

An agent that practices your task for a few minutes before you trust it to act

What we provide

On-demand practice environments

One set of environments serves all three, **and every model upgrade needs new tasks**

Our edge

Quant-grade rigor, **plus Asia's professional talent**

01

Quant and statistical research

Hidden tests, out-of-sample evaluation and adversarial behavior are everyday practice in quant research. We build every environment the same way.

Xitadel separates models from 24 to 77, against a human best of 80

02

Asian professional talent and workflows

Direct access to professionals and how they actually work lets us build vertical environments fast, without lowering the technical bar.

Seven environments shipped in our first 14 days

Practitioners in our R&D come from top trading firms, including



200K+

Reachable experts

7,000+

Experts on the waitlist

Write, score, review

What experts do in our environments

Logos show where practitioners work; no endorsement implied.

Talent network

Students and alumni of 21 top universities



21

Universities in our student and alumni network

From entry-level roles to frontier research

business@simreal.co · simreal.com.cn

Selected universities shown; no endorsement implied.