

Every industry's best agents will be trained in our worlds

An AI-native neolab building RL environments, verifiers and training data for agents that improve themselves

Founded

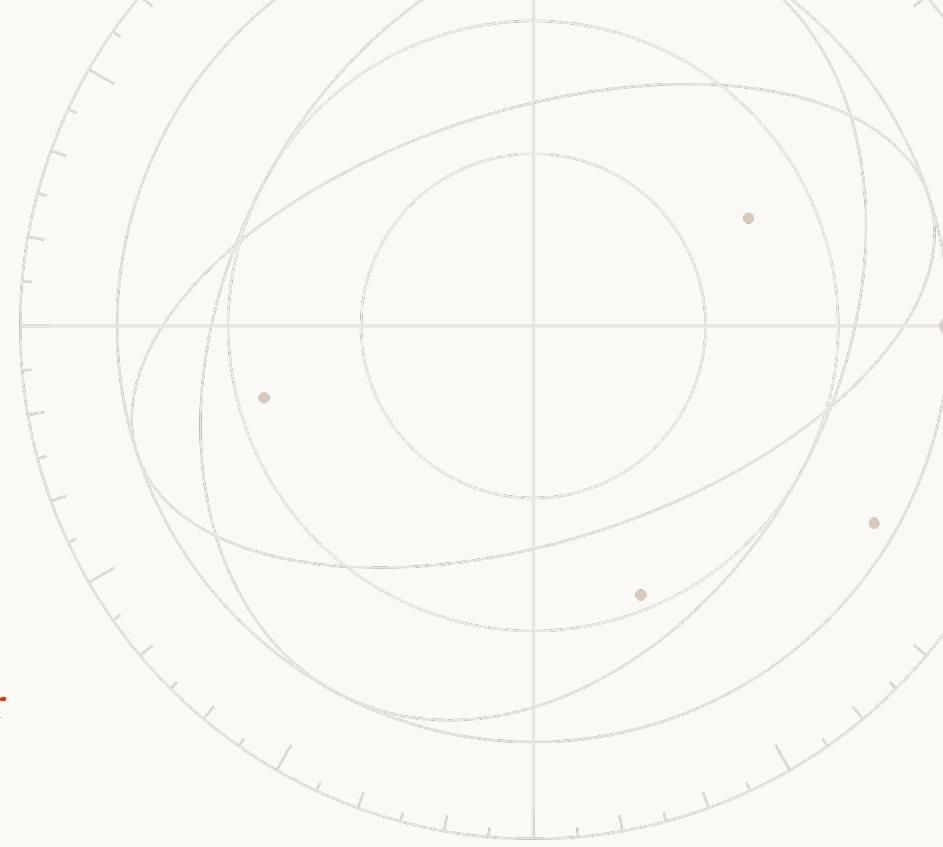
September 10, 2026

Raising

\$6M seed round

Investor presentation

October 2026



Overview

An AI-native neolab building RSI engines: **the environments, verifiers and data that let agents improve themselves**

Revenue

\$3M ARR

Reached 24 days after founding
No outside capital

Product

7 environments

Trading, ML research, math, reasoning,
forecasting, software and accounting

Proof

+12%

A model trained in our trading environment
did better on trading days it had never seen

Open source

602 GitHub stars

Across 5 open benchmarks,
all shipped in our first 14 days

Team

Quant founders, all 21

Cambridge, LSE and Duke mathematics
Jane Street, Citadel, Optiver, Millennium

Market

\$8.5B a year

Spent today on training data
and RL environments

Team

Three founders, all 21. Together they turned down \$0.9M+ in first-year pay to build SimReal



Charles CEO · 21

First employee at United Stables, where \$U grew from zero to \$1.4B in a year and listed on Binance within a month

Led institutional partnerships with YZi Labs, SIG, DRW and Wintermute

Only intern on HSBC's HKD stablecoin project, working on HKMA compliance

Hedge fund intern, Citadel London

LSE, Mathematics
USAMO qualifier

Henry CTO · 21

Cambridge Mathematics, First Class; scholar of St John's College

Youngest researcher at the Cambridge AI research center, under Prof. Po-Ling Loh (IMS Fellow)

Ex-quant at Jane Street, Citadel and Optiver

Designed all 5 open benchmarks

Top 30 worldwide, Cambridge maths competition
British Physics Olympiad Top Gold

Amaris COO · 21

Data scientist, Millennium Hong Kong

First graduate hire on Millennium's alternative-data team

Co-organized Plug and Play's first Hong Kong event, with HKSTP (200+ attendees)

Duke, Mathematics and Statistics

Problem

Agents can reason. **They still fail at real work**

	Looks right		What actually happens
Trading	Wins on past data	→	Loses money live
Software	All tests pass	→	Breaks in production
Accounting	Books look closed	→	Month-end doesn't reconcile
Forecasting	Analysis sounds right	→	Outcome proves it wrong


6/32

In a live trading contest (Alpha Arena), frontier models made money in only 6 of 32 runs

The gap: **agents have nowhere realistic to practice, and no source of truth to learn from**

Solution

A loop that turns real outcomes into better agents, **with AI inside it from day one**



What's different

Today

Humans hand-write every task, label and rubric

SimReal

Agents help write new tasks and data, find failure modes and improve the environments they train in. Next, they will write the verifiers too

Proof point

+12%

on trading days it never saw,
after training in our loop

+12% is the best of several independent replications, measured on trading days held out from training.

Product

Seven environments in 14 days. Each is scored by something the agent can't argue with

602

GitHub stars across
5 open benchmarks

As of October 4, 2026

Give the agent a world to practice in.


Make the reward come from a market, a proof, hidden labels, an executable system or a real-world outcome.

Domain	Environment	Scored by	Stars
Trading	Xitadel	Profit on unseen market days, vs. the best human	102
ML research	SimReal-MLBench	The official Kaggle leaderboard, on hidden labels	302
Math	MathmoBench	Whether the proof is correct	101
Reasoning	Puzzle Benchmark	Hidden answers to 749 human-written puzzles	70
Forecasting	FuturePredict	The real outcome, after the deadline	27
Software	SWE-Forward	Hidden tests the code must pass	Private
Accounting	Month-End Close	Whether the books reconcile	Private

SWE-Forward and Month-End Close are not open-sourced.

How it works


Three environments. Each scores the AI against something real



Best human: on each task, the best real strategy from IMC Prosperity, a global trading competition. FuturePredict baseline: a reference forecast, also locked in advance.

Traction

Founded September 10, 2026. **\$3M ARR in 24 days, with no outside capital**



Days counted from founding; as of October 4, 2026

Infrastructure figures come from separate internal tests and should not be added together.

Why now

Agents are starting to improve themselves. **Environments are now the bottleneck**

What changed in 2026		What we have built
AI is doing AI research OpenAI: agents do 3.1 days of work for every researcher-day OpenAI, Sep 2026	→	SimReal-MLBench: an exam for AI research 302 GitHub stars
AI is improving itself An agent rewrote its own code and improved 7 times in 8 days Weco AIDE², Sep 2026	→	RSI working in trading Up to +12% on unseen trading days
Environments are the bottleneck When the environment gets noisy, top agents drop from 83.9% to 57.6% Breaking the Environment Wall, Sep 2026	→	Seven environments in 14 days 200K+ experts to build more

Our position

Environments gate RSI; they are our core research. **Most neolabs have no revenue; we hit \$3M ARR in 24 days**

Sources: OpenAI (Sep 6, 2026); Weco AI, arXiv 2609.26457; Breaking the Environment Wall, arXiv 2609.29773; Deedy Das neolab list (May 2026). +12% is the best of several replications.

Market

\$8.5B a year today. **One vendor, Mercor, grew revenue 27x in 16 months**

Today: training data and RL environments

\$8.5B / year

Combined revenue of 50+ vendors

2030: the AI economy

→ **\$700B / year**

If AI becomes a \$7T economy and 10% goes to training

27x

Mercor run-rate revenue:
\$75M to \$2B in 16 months

→ **Demand for expert data is exploding**

10x

Mercor valuation:
\$2B to \$20B in 17 months (in talks)

→ **Investors are paying up**

18x

Snorkel AI's new data business:
growth in one year

→ **Demand is shifting to experts and environments**

Sources: Deedy Das market map (Jul 2026); Mercor (TechCrunch, Dealroom); Snorkel AI (Sep 2026). The 2030 figure is our estimate.

Business model

We sell environments, data and evals. **Frontier labs pay today; enterprise and personal agents next**

Now

Frontier AI labs

They buy

Environments, verified tasks and evals to post-train their models

How we charge

Per environment and per task, on quarterly contracts

Estimated price

\$300K–\$1M+

Per lab, per quarter

Market rates: \$200–\$2,000 a task; about \$300K for a complex environment

Next

Enterprise agents

They buy

Custom environments and tests that prove an agent can do the job

How we charge

A paid pilot, then an annual license

Estimated price

\$200K–\$800K

Per customer, per year

Pilot: \$50K–\$150K over 8–12 weeks

Later

Personal agents

They buy

Short practice runs that fit an agent to one person's habits before it acts for them

How we charge

A monthly subscription, direct or through agent apps

Estimated price

\$5–\$10

Per user, per month

For reference, consumer AI plans cost about \$20 a month

One set of environments serves all three, **and every new model needs new tasks**

Estimates, not signed prices. Lab prices are market rates from Epoch AI interviews with environment builders and AI labs (January 2026). Enterprise and personal prices are our assumptions.

Why us

Data vendors sell expert hours. **We build assets that compound with every customer**

	Expert-data vendors	SimReal
What compounds	Expert hours. Each project starts from scratch.	Proprietary environments, verifiers and maps of where models fail, plus the models we train in them. Every customer and loop adds to them.
Cost curve	Rises with headcount	Falls with scale: once a verifier exists, new data is generated and checked automatically
Proof of quality	Spot checks and reviewer judgment	Hidden tests, out-of-sample results, hash-locked published results and measured training gains (+12%)
Velocity	A new domain means new recruiting	7 environments in 14 days; \$3M ARR in 24 days
Network	Large expert pools, paid by the hour	Practitioners from top trading firms and 200K+ reachable experts, whose work our verifiers check

In short

They sell hours; we sell verified gains. **Their costs rise with scale; ours fall**

Network

Experts from the most competitive entry-level jobs and from research, **the work agents are now built to do**

Practitioners in our R&D come from top trading firms and markets, including



The most competitive entry-level jobs

Trading, finance, software and more

→ **Data for agents that streamline everyday work**

Research

AI, mathematics and quantitative research

→ **Data for agents that automate research**

200K+

Reachable experts

7,000+

Experts on the waitlist

21

Top universities

Students and alumni of



Logos show where practitioners work or study; no endorsement implied.

The round

Raising a \$6M seed round **to scale environments, RSI and delivery**

Raising

\$6M

Seed round

Three directions, each with a milestone →

01

Environments and data

Take the trading playbook to more expert domains

Where the money goes

Environment engineers ·
Expert network

Milestone

10+ expert-domain environments

02

RSI

Scale self-improvement to larger models and more domains

Where the money goes

Compute · Researchers

Milestone

RSI on open-weight large models

03

Delivery

Turn early demand into lasting contracts

Where the money goes

Delivery team · Sales

Milestone

20 customers

business@simreal.co · simreal.com.cn