



Every industry's best agents will be trained in our worlds

An AI-native neolab building environments, verifiers
and data for self-improving AI

Founded


September 10, 2026

Raising

\$6M seed round

Contact

business@simreal.co



Overview

An AI-native neolab building **the environments, verifiers and data that let agents improve themselves**

Revenue

\$3M run-rate

Gross, from our expert-data business that predates SimReal; no outside capital

Flagship

Trading environment

Market making, options, baskets; scored on P&L on a withheld day

Pipeline

2 frontier labs

In direct talks; today we sell through data vendors that serve these labs

Team

Three founders, all 21

Cambridge, LSE and Duke mathematics
Interned at Jane Street, Citadel, Optiver

Open source

600+ GitHub stars

Across our open benchmarks, as of October 2026

Market

\$8.5B a year

Spent today on training data and RL environments

How it fits

Our data business funds the environments. The environments are built to make that data verifiable and cheaper to produce. Training open models in them closes the loop.

Team

Three 21-year-old founders. Together they turned down \$800K+ in first-year pay to build SimReal

Where we have worked or interned



Charles CEO · 21

Employee #1 at United Stables: \$U listed on Binance within a month of launch and grew to \$1.4B+ in circulation

Led institutional partnerships with YZi Labs, SIG, DRW and Wintermute

Only intern on HSBC's HKD stablecoin project, working on HKMA compliance

Hedge fund intern, Citadel

LSE, Mathematics
USAMO qualifier

Henry CTO · 21

Cambridge Mathematics, St John's College, First Class with scholarship

Research with Prof. Po-Ling Loh at Cambridge

Quant internships: Jane Street, Citadel, Optiver

Top 30 worldwide on STEP, Cambridge's hardest maths entrance exam
British Physics Olympiad Top Gold

Amaris COO · 21

Data scientist, Millennium Hong Kong

First graduate hire on Millennium's alternative-data team

Co-organized Plug and Play's first Hong Kong event, with HKSTP (200+ attendees)

Duke, Mathematics and Statistics

Problem

Agents can reason. **They still fail at real work**

	Looks right		What actually happens
Trading	Profitable on the data it was tuned on	→	Loses money on new days
Software	All tests pass	→	Breaks in production
Accounting	Books look closed	→	Month-end doesn't reconcile
Forecasting	Analysis sounds right	→	Outcome proves it wrong

6/32

In a live trading contest (Nof1's Alpha Arena), frontier models made money in only 6 of 32 runs

The gap: **too few realistic places to practice, and too little ground truth to learn from**

Source: Nof1, Alpha Arena (May 2026).

Why now

Agents are starting to improve themselves. **Environments are now the bottleneck**

What changed in 2026		What we have built
AI is doing AI research OpenAI: agents do 3.1 days of work for every researcher-day OpenAI, Sep 2026	→	SimReal-MLBench: an exam for ML research agents 60 past Kaggle competitions, compared with humans
AI is improving itself An agent rewrote its own code and improved 7 times in 8 days Weco AIDE², Sep 2026	→	Training loop running in trading Train on past days; score on withheld trading days
Environments are the bottleneck When the environment gets noisy, top agents drop from 83.9% to 57.6% Breaking the Environment Wall, Sep 2026	→	Seven environments in 14 days 200K+ students and alumni to build more


Our position

Environments gate self-improvement. **Most neolabs have no revenue; we already have a \$3M run-rate**

Sources: OpenAI (Sep 6, 2026); Weco AI, arXiv 2609.26457; Breaking the Environment Wall, arXiv 2609.29773; Deedy Das neolab list (May 2026).

Solution

A loop that turns real outcomes into better agents, **with AI inside it from day one**



Loop figures come from separate internal tests and should not be added together.

What's different

Today

Tasks, labels and rubrics are mostly written by hand

SimReal

Agents already help us write tasks and find where models fail. In development: agents that write the verifiers and improve the environments themselves.

Training infrastructure: more per unit of compute

1/4

the time per
practice run

+64%

scored attempts,
same compute

-1/3

time spent saving
checkpoints

Product

Trading is the flagship environment, **scored by P&L: something the agent can't argue with**

FLAGSHIP

•

TRADING

Xitadel

open source

An agent researches a market, codes a strategy and is scored on a trading day withheld from it.

Seven published tasks, five types

Single-product market making

Options vs. underlying

Baskets vs. components

Conversion with tariffs and storage

Multi-asset, disclosed counterparties

Scored by

P&L on a withheld trading day, penalized for drawdown and low Sharpe, against the strongest human strategy on each task (set at 80 of 100)

Supporting environments

ML research

SimReal-MLBench

60 past Kaggle competitions, compared with the final leaderboard

Forecasting

FuturePredict

Questions with no answer yet, scored on the resolved outcome against a sealed baseline

Also built

Math proofs (MathmoBench) · Reasoning (Puzzle Benchmark)

Accounting (Month-End Close)

Software (SWE-Forward): must survive the next real release

600+ GitHub stars across our open benchmarks

SWE-Forward and Month-End Close are private. Human reference: the strongest real strategy on each task from IMC Prosperity, a trading competition. GitHub stars as of October 2026.

Inside the trading environment

Research, code, then a private held-out day. **Scored on P&L against the best human strategy**

01

Research

Historical order books and trades for 2-8 training days, with position limits and product rules



02

Code

The agent backtests locally and submits a Python strategy: `Trader.run(state) → orders`



03

Held-out replay

Run step by step on a private test day. At each step the strategy sees only the order books, reported trades and its own position

Published results · score out of 100; human reference = 80

Best human	80
GPT 6	77.3
GLM 5.3	30.1
Kimi K3	27.7
DeepSeek V4 Pro	24.6

Human reference: the strongest real strategy on each task from IMC Prosperity, a trading competition. Seven published tasks across five task types.

Score and controls

Score	max(0, P&L - 0.1 × max drawdown) × Sharpe factor (0.8-1.0) Rescaled: best human = 80, twice it = 100 (the cap)
Selection	Tasks fixed before any model was scored; reserve test days replace any that become public
Audit	Full results hash-committed (SHA-256); the training reward is separate from this score

Traction

\$3M revenue run-rate, before any outside capital. **Sold through data vendors that serve frontier labs**

Revenue

\$3M

Gross revenue run-rate

From our expert-data business, which began before SimReal was incorporated

Sold through data vendors that serve frontier labs; orders are placed on demand

Direct lab contracts are in talks

Since founding on September 10, 2026

Day 14	7 environments shipped Trading, ML research, forecasting, math, reasoning, software, accounting
Training	Training loop running end to end Models train on past trading days and are scored on a withheld day
Data	10+ types of expert data delivered Through our expert-data business
Open source	600+ GitHub stars Across our open benchmarks, as of October 2026
Pipeline	2 frontier labs in direct talks For environment and data contracts

Run-rate is gross: it includes what we pay the experts who produce the data.

Why us

Data vendors sell expert hours. **We build assets that compound with every customer**

Today we sell expert data through data vendors that serve frontier labs. The environments and verifiers we build ourselves stay with us.

	Expert-data vendors	SimReal
What compounds	Mostly expert hours	Environments, verifiers and maps of where models fail, plus the models we train in them
Cost curve	Rises with headcount	Built to fall with scale: once a verifier exists, new data is checked automatically
Proof of quality	Mostly human review	Hidden tests, out-of-sample scoring and hash-committed (SHA-256) results
Velocity	A new domain means new recruiting	7 environments in 14 days

Landscape

Expert-data vendors	Mercor, Surge AI, Scale AI	Mostly sell expert hours; some now build environments
Environment startups	Mechanize, Halluminate, AfterQuery, UniPat	Build environments for AI labs
SimReal	Expert data and environments	Environments scored by real outcomes; data that already earns revenue

In short

They sell hours; we sell verified gains. **Their costs rise with scale; ours fall**

Network

Students, alumni and practitioners from top entry-level jobs and research, **the work agents are now built to do**

Practitioners in our R&D have worked at top trading firms and markets, including



The most competitive entry-level jobs

Trading, finance, software and more

→ **Data for agents that streamline everyday work**

Research

AI, mathematics and quantitative research

→ **Data for agents that automate research**

200K+

**Students and alumni
we can reach**

7,000+

**Signed up for
paid data work**

21

Top universities

Students and alumni of



Logos show where practitioners have worked or studied; no endorsement implied.

Market

\$8.5B a year today. **One vendor, Mercor, grew its gross run-rate 27x in 16 months**

Today: training data and RL environments

\$8.5B / year

Combined run-rates of 50+ data and environment vendors

Per-account estimate: one frontier lab

\$1.2M–\$4M+ / year

Our estimate at market rates for environments and tasks

27x

Mercor gross run-rate:
\$75M to \$2B in 16 months

→ **Demand for expert data is growing fast**

18x

Snorkel AI's new data business:
growth in one year

→ **Demand is shifting to experts and environments**

Sources: Deedy Das market map (Jul 2026); Mercor (TechCrunch, Dealroom); Snorkel AI (Sep 2026); market rates from Epoch AI (Jan 2026). Vendor run-rates are partly gross.

Business model

We sell expert data today, through vendors that serve frontier labs. **Environments, evals and enterprise next**

Now

Frontier AI labs

They buy

Expert data and verified tasks today;
environments and evals in direct talks

How we charge

Today per task, through data vendors;
next, direct quarterly contracts

Estimated price

\$300K–\$1M+

Per lab, per quarter, direct

Market rates: \$200–\$2,000 a task; about \$300K for a complex environment

Next

Enterprise agents

They buy

Custom environments and tests that
prove an agent can do the job

How we charge

A paid pilot, then an annual license

Estimated price

\$200K–\$800K

Per customer, per year

Pilot: \$50K–\$150K over 8–12 weeks

Later

Personal agents

They buy

Fit an agent to one person's habits and
strategies before it acts for them

How we charge

A usage fee, paid by the company that
runs the agent

Pricing

Per user

Per active user, per month

Paid by the agent company, not the consumer

The same environments and verifiers carry across all three, **and every new model needs new tasks**

Estimates, not signed prices. Lab prices are market rates from Epoch AI interviews with environment builders and AI labs (January 2026). Enterprise prices are our assumptions.

The round

Raising a \$6M seed round **to scale environments, self-improvement research and delivery**

Raising

\$6M

Seed round

Three directions, each with a milestone →

01

Environments & data

Take the trading playbook to more expert domains

Where the money goes
Engineers · Expert network

Milestone

10+ expert-domain environments

02

Self-improvement

Scale the training loop to larger models and domains

Where the money goes
Compute · Researchers

Milestone

Training loop on open-weight large models

03

Delivery

Turn on-demand orders into direct, recurring contracts

Where the money goes
Delivery team · Sales

Milestone

20 customers

business@simreal.co