simreal

Every industry's best agents will be trained in our worlds

An AI-native neolab building environments, verifiers and data for self-improving AI

Founded	Raising	Investor presentation
September 10, 2026	Seed round	October 2026

Overview

An AI-native neolab building RSI engines: the environments, verifiers and data that let agents improve themselves

Revenue

$3M ARR

Reached 24 days after founding
No outside capital

Product

7 environments

Trading, ML research, math, reasoning, forecasting, software and accounting

Proof

+12%

A model trained in our trading environment did better on trading days it had never seen

Open source

672 GitHub stars

Across 5 open benchmarks, all shipped in our first 14 days

Team

Quant founders, all 21

Cambridge, LSE and Duke mathematics
Jane Street, Citadel, Optiver, Millennium

Market

$8.5B a year

Spent today on training data and RL environments

Team

Four 21-year-old founders. Together they turned down $1M+ in first-year pay to build SimReal

Jane Street · CITADEL · millennium · Optiver · University of Cambridge

Charles CEO · 21

Employee #1 at United Stables, building it from scratch to $1.5B+. Listed on Binance within a month.

Led institutional partnerships with YZi Labs, SIG, DRW and Wintermute

Only intern on HSBC's HKD stablecoin project, working on HKMA compliance

Hedge fund intern, Citadel

LSE, Mathematics
USAMO qualifier

Henry CTO · 21

Cambridge Mathematics, St John's College, First Class with Scholarship

Youngest researcher at the Cambridge AI research center, under Prof. Po-Ling Loh (IMS Fellow)

Ex-quant at Jane Street, Citadel and Optiver

Top 30 worldwide, Cambridge STEP maths competition (~IMO Medalist level)
British Physics Olympiad Top Gold

Amaris COO · 21

Data scientist, Millennium Hong Kong

First graduate hire on Millennium's alternative-data team

Co-organized Plug and Play's first Hong Kong event, with HKSTP (200+ attendees)

Duke, Mathematics and Statistics

James CPO · 21

Quant trading at Optiver

Asset management at Fidelity

Equity research at Jefferies

LSE, Mathematics
AMC: global top 1%
British Mathematical Olympiad: Gold

Problem

Agents can reason. They still fail at real work

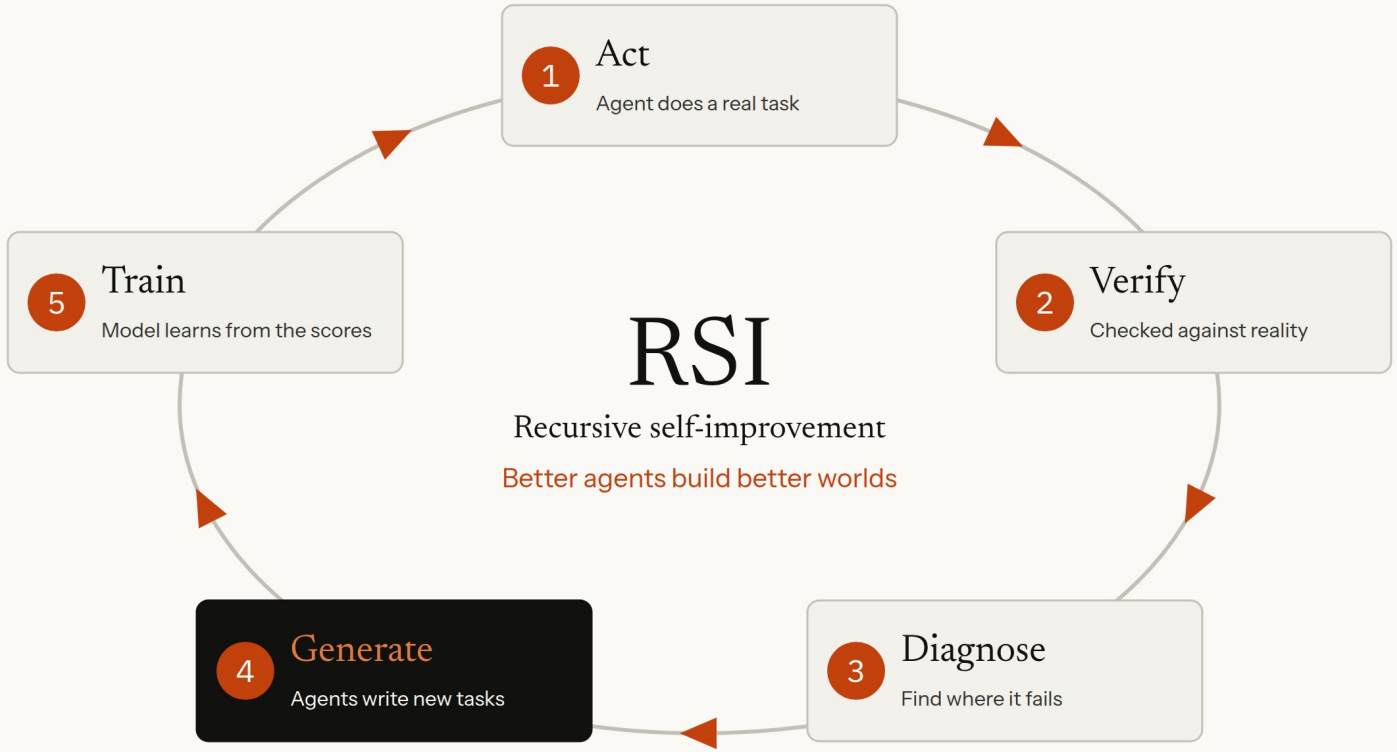
	Looks right		What actually happens
Trading	Wins on past data	→	Loses money live
Software	All tests pass	→	Breaks in production
Accounting	Books look closed	→	Month-end doesn't reconcile
Forecasting	Analysis sounds right	→	Outcome proves it wrong

6/32 In a live trading contest (Alpha Arena), frontier models made money in only 6 of 32 runs

The gap: agents have nowhere realistic to practice, and no source of truth to learn from

Solution

A loop that turns real outcomes into better agents, with AI inside it from day one



What's different

Today

Humans hand-write every task, label and rubric

SimReal

Agents help write new tasks and data, find failure modes and improve the environments they train in. Next, they will write the verifiers too. AI native from the start.

Proof point

+12% on trading days it never saw, after training in our loop

+12% is the best of several independent replications, measured on trading days held out from training.

Product

Seven environments in 14 days. Each is scored by something the agent can't argue with

672

GitHub stars across 5 open benchmarks

As of October 4, 2026

Give the agent a world to practice in.

Make the reward come from a market, a proof, hidden labels, an executable system or a real-world outcome.

Domain	Environment	Scored by	Stars
Trading	Xitadel	Profit on unseen market days, vs. the best human	101
ML research	SimReal-MLBench	The official Kaggle leaderboard, on hidden labels	301
Math	MathmoBench	Whether the proof is correct	101
Reasoning	Puzzle Benchmark	Hidden answers to 749 human-written puzzles	100
Forecasting	FuturePredict	The real outcome, after the deadline	27
Software	SWE-Forward	Hidden tests the code must pass	Private
Accounting	Month-End Close	Whether the books reconcile	Private

SWE-Forward and Month-End Close are not open-sourced.

How it works

Three environments. Each scores the AI against something real

	1 Given	2 The AI	3 Scored against

Trading
Xitadel
Published results
	7 trading tasks
	Spot ×2 · Options ×2 · Basket · Conversion · Multi-asset
		Writes a trading strategy
		`class Trader:`
		`  def run(self, state):`
		`  return orders`
		Tested on a market day it never saw
			The best human on each task
			Best human 80
			GPT 6 77.3
			GLM 5.3 30.1
			Kimi K3 27.7
			DeepSeek V4 Pro 24.6

AI research SimReal-MLBench Baseline result, not an AI agent	**60 past Kaggle competitions** For example — Time limit House prices — 6 h Dogs vs. cats — 12 h Bird calls — 24 h	**Runs the whole project** Data → Features → Training → Submit Best of up to 2 submissions	**The final Kaggle leaderboard** Bike Sharing Demand competition A baseline beat **84%** of 3,242 human teams Last place — First place
Forecasting FuturePredict Illustrative numbers	**A question with no answer yet** USGS earthquake data How many magnitude 5+ earthquakes worldwide tomorrow? 0 · 1 · 2+	**Gives its odds** 0 10% 1 30% 2+ **60%** Locked before anyone knows the answer	**The real outcome: 2+ quakes** illustrative Error (lower is better) Baseline 0.33 AI 0.13 Reward 0.20

Best human: on each task, the best real strategy from IMC Prosperity, a global trading competition. FuturePredict baseline: a reference forecast, also locked in advance.

Traction

Founded September 10, 2026. $3M ARR in 24 days, with no outside capital

Day 14

7

Environments shipped

Trading training loop working

Week 3

10+

Types of data delivered

Day 24

$3M

ARR

Day 31

672

GitHub stars

Across 5 open benchmarks

Days counted from founding; as of October 4, 2026

Customers

2

Frontier labs in talks

`Training infrastructure: more training per unit of compute`

1/4

Each practice run takes a quarter of the time

+64%

More scored attempts on the same compute

−1/3

Less time spent saving model checkpoints

Infrastructure figures come from separate internal tests and should not be added together.

Why now

Agents are starting to improve themselves. Environments are now the bottleneck

What changed in 2026	What we have built
AI is doing AI research OpenAI: agents do 3.1 days of work for every researcher-day OpenAI, Sep 2026	**SimReal-MLBench: an exam for AI research** 302 GitHub stars
AI is improving itself An agent rewrote its own code and improved 7 times in 8 days Weco AIDE², Sep 2026	**RSI working in trading** Up to +12% on unseen trading days
Environments are the bottleneck When the environment gets noisy, top agents drop from 83.9% to 57.6% Breaking the Environment Wall, Sep 2026	**Seven environments in 14 days** 200K+ experts to build more

Our position Environments gate RSI; they are our core research. **Most neolabs have no revenue; we hit $3M ARR in 24 days**

Sources: OpenAI (Sep 6, 2026); Weco AI, arXiv 2609.26457; Breaking the Environment Wall, arXiv 2609.29773; Deedy Das neolab list (May 2026). +12% is the best of several replications.

Market

$8.5B a year today. One vendor, Mercor, grew revenue 27x in 16 months

Today: training data and RL environments

$8.5B / year

Combined revenue of 50+ vendors

→

2030: the AI economy

$700B / year

If AI becomes a $7T economy and 10% goes to training

27x

Mercor run-rate revenue:
$75M to $2B in 16 months

→ **Demand for expert data is exploding**

10x

Mercor valuation:
$2B to $20B in 17 months (in talks)

→ **Investors are paying up**

18x

Snorkel AI's new data business:
growth in one year

→ **Demand is shifting to experts and environments**

Sources: Deedy Das market map (Jul 2026); Mercor (TechCrunch, Dealroom); Snorkel AI (Sep 2026). The 2030 figure is our estimate.

Business model

We sell environments, data and evals. Frontier labs pay today; enterprise and personal agents next

Now

Frontier AI labs

They buy

Environments, verified tasks and evals to post-train their models

How we charge

Per environment and per task, on quarterly contracts

Estimated price

$300K–$1M+

Per lab, per quarter

Market rates: $200–$2,000 a task; about $300K for a complex environment

Next

Enterprise agents

They buy

Custom environments and tests that prove an agent can do the job

How we charge

A paid pilot, then an annual license

Estimated price

$200K–$800K

Per customer, per year

Pilot: $50K–$150K over 8–12 weeks

Later

Personal agents

They buy

Fit an agent to one person's habits and strategies before it acts for them

How we charge

A monthly subscription, direct or through agent apps

Estimated price

$5–$10

Per user, per month

For reference, consumer AI plans cost about $20 a month

One set of environments serves all three, and every new model needs new tasks

Estimates, not signed prices. Lab prices are market rates from Epoch AI interviews with environment builders and AI labs (January 2026). Enterprise and personal prices are our assumptions.

Why us

Data vendors sell expert hours. We build assets that compound with every customer

	Expert-data vendors	SimReal
What compounds	Expert hours. Each project starts from scratch.	**Proprietary environments, verifiers and maps of where models fail, plus the models we train in them. Every customer and loop adds to them.**
Cost curve	Rises with headcount	**Falls with scale: once a verifier exists, new data is generated and checked automatically**
Proof of quality	Spot checks and reviewer judgment	**Hidden tests, out-of-sample results, hash-locked published results and measured training gains (+12%)**
Velocity	**A new domain means new recruiting**	**7 environments in 14 days; $3M ARR in 24 days.**
Network	Large expert pools, paid by the hour	**Practitioners from top trading firms and 200K+ reachable experts, whose work our verifiers check**

In short — They sell hours; we sell verified gains. **Their costs rise with scale; ours fall**

Network

Experts from the most competitive entry-level jobs and from research, the work agents are now built to do

Practitioners in our R&D come from top trading firms and markets, including

Jane Street · imc · CITADEL | Securities · Optiver · Polymarket

The most competitive entry-level jobs

Trading, finance, software and more

→ **Data for agents that streamline everyday work**

Research

AI, mathematics and quantitative research

→ **Data for agents that automate research**

200K+	7,000+	21
Reachable experts	**Experts on the waitlist**	**Top universities**

Students and alumni of

Harvard University · Stanford University · MIT · University of Oxford · University of Cambridge · Princeton University · Yale University · Berkeley University of California · 清華大學 · Columbia University · The University of Chicago · Duke University

Logos show where practitioners work or study; no endorsement implied.